

#4 Terraform Associate Certification 003 - Results

[Back to result overview](#)

Attempt 1

All domains

57 all

38 correct

19 incorrect

0 skipped

0 marked

[Collapse all questions](#)

Question 1 **Correct**

terraform validate confirms the syntax of Terraform files.

Your answer is correct

True

Explanation

True. The `terraform validate` command is used to check the syntax of Terraform configuration files. It ensures that the files are correctly formatted and do not contain any syntax errors that could cause issues during deployment.

False

Explanation

False. The `terraform validate` command is specifically designed to confirm the syntax of Terraform files. It does not perform any other functions such as executing the code or validating the logic of the configuration.

Overall explanation

Correct Answer: True

Explanation:

The `terraform validate` command checks the syntax and internal consistency of your Terraform configuration files (such as `.tf` files). It ensures that the files are correctly formatted and that the configurations do not contain any obvious errors. However, it does not check the validity of the infrastructure in relation to the actual cloud provider or environment—it only verifies that the configuration is syntactically correct and can be processed by Terraform.

This command does not apply any changes to the infrastructure, nor does it interact with the provider API.

Resources

[Command: validate](#)

Question 2 **Incorrect**

Which of the following module source paths does not specify a remote module?

Correct answer

source = `"/modules/consul"`

Explanation

The source path `"/modules/consul"` specifies a local module, not a remote module. It indicates that the module is located in the local directory structure, rather than being fetched from a remote repository like GitHub or Terraform Registry.

source = `"hashicorp/consul/aws"`

Explanation

The source path `"hashicorp/consul/aws"` specifies a remote module from the Terraform Registry. This source path indicates that the module should be fetched from the Terraform Registry under the namespace `"hashicorp"` and the module name `"consul/aws"`.

source = `"github.com/hashicorp/example"`

Explanation

The source path `"github.com/hashicorp/example"` specifies a remote module from a GitHub repository. This source path indicates that the module should be fetched from a specific GitHub repository using the HTTPS protocol.

Your answer is incorrect

source = `"git@github.com:hashicorp/example.git"`

Explanation

The source path "git@github.com:hashicorp/example.git" specifies a remote module from a Git repository hosted on GitHub. This source path indicates that the module should be fetched from a specific Git repository using the SSH protocol.

Overall explanation

The correct answer is: source = "./modules/consul"

Explanation:

A module source path specifies where Terraform should find the module.

- **"./modules/consul"**: This is a relative path to a local module on disk. It does not specify a remote module because the module resides locally in the directory structure.
- **"git@github.com:hashicorp/example.git"**: This specifies a remote module sourced from a Git repository using SSH.
- **"github.com/hashicorp/example"**: This specifies a remote module sourced from GitHub using the HTTPS protocol.
- **"hashicorp/consul/aws"**: This specifies a remote module from the **Terraform Registry**, which is also considered remote.

Key takeaway:

Local paths (e.g., `./modules/consul`) are not remote sources, whereas Git repositories, Terraform Registry modules, and other URLs are considered remote sources.

Resources

[Module Sources](#)

Question 3Correct

You are making changes to existing Terraform code to add some new infrastructure.

When is the best time to run `terraform validate`?

After you run terraform plan so you can validate that your state file is consistent with your infrastructure

Explanation

Running terraform validate after terraform plan may not be the best approach because terraform validate checks the syntax of your Terraform code before creating an execution plan. It is more efficient to catch any syntax errors early on before generating a plan.

Your answer is correct

Before you run terraform plan so you can validate your code syntax

Explanation

The best time to run terraform validate is before running terraform plan. This allows you to validate your code syntax and catch any errors or issues before proceeding to create an execution plan. By validating the code first, you can ensure that your Terraform configuration is correct and ready for planning.

After you run terraform apply so you can validate that your infrastructure is reflected in your code

Explanation

Running terraform validate after terraform apply is not recommended because terraform validate is intended to check the syntax of your Terraform code before creating any infrastructure. Validating after applying changes may not be as effective in catching potential errors in your code.

Before you run terraform apply so you can validate your infrastructure changes

Explanation

Running terraform validate before terraform apply is also a valid approach, as it allows you to validate your infrastructure changes before applying them. However, running it before terraform plan is more efficient as it helps ensure that your code syntax is correct before generating a plan.

Overall explanation

Correct Answer: Before you run terraform plan so you can validate your code syntax

Explanation:

The `terraform validate` command is used to check the syntax and internal consistency of your Terraform configuration files without accessing any external services, such as the remote state or cloud provider APIs. Running it before `terraform plan` ensures that your code is free of syntax errors and that all necessary configurations are properly defined. This is an essential step to catch basic issues early in the workflow.

Why other answers are incorrect:

- **After you run `terraform plan` so you can validate that your state file is consistent with your infrastructure:** `terraform validate` does not check state files or infrastructure consistency; it only validates configuration syntax.
- **Before you run `terraform apply` so you can validate your infrastructure changes:** While it's fine to validate code before applying changes, the best time to run `terraform validate` is before `terraform plan`, as `plan` relies on correct configuration syntax to generate a plan.
- **After you run `terraform apply` so you can validate that your infrastructure is reflected in your code:** `terraform validate` does not check the actual infrastructure or its consistency with the configuration. This would be done using commands like `terraform plan` or `terraform refresh`.

Resources

[Terraform validate](#)

Question 4Correct

In Terraform HCL, an object type of `object({ name=string, age=number })` would match this value:

```
{  
  
  name = "John"  
  
  age = fifty two  
  
}
```

Explanation

Choice A is incorrect because the value for the "age" attribute is not a valid number. In Terraform HCL, the age attribute should be assigned a number, not a string like "fifty two".

Your answer is correct

```
{  
  
  name = "John"  
  
  age = 52  
  
}
```

Explanation

Choice B is correct because it matches the object type specified in the Terraform HCL definition. The "name" attribute is assigned a string value "John", and the "age" attribute is assigned a number value 52, which aligns with the object type definition.

```
{  
  
  name = John  
  
  age = fifty two  
  
}
```

Explanation

Choice C is incorrect because the value for the "name" attribute is not enclosed in double quotes, making it an invalid string in Terraform HCL. Additionally, the value for the "age" attribute is not a valid number.

```
{  
  
  name = John  
  
  age = 52
```

```
}
```

Explanation

Choice D is incorrect because the value for the "name" attribute is not enclosed in double quotes, making it an invalid string in Terraform HCL. The "age" attribute has a valid number value, but the format for the "name" attribute is incorrect.

Overall explanation

Correct Answer:

```
{ name = "John" age = 52 }
```

Explanation:

In Terraform's HCL (HashiCorp Configuration Language), an object type like `object({ name=string, age=number })` requires the values for the keys to match their respective types:

- `name` must be a string.
- `age` must be a number.

The correct value provided is `{ name = "John" age = 52 }` because:

1. The `name` key has a string value `"John"`.
2. The `age` key has a numeric value `52`.

Why other options are incorrect:

- `{ name = "John" age = fifty two }`: `fifty two` is not a number but a string, which does not match the `number` type for `age`.
- `{ name = John age = fifty two }`: Missing quotation marks around `John` makes it invalid for the `string` type, and `fifty two` is not a number.
- `{ name = John age = 52 }`: Missing quotation marks around `John` violates the `string` type requirement for `name`.

Question 5Correct

What is a key benefit of the Terraform state file?

A state file can be used to schedule recurring infrastructure tasks

Explanation

A state file in Terraform is not used to schedule recurring infrastructure tasks. It is primarily used to store the current state of the infrastructure managed by Terraform.

A state file represents a source of truth for resources provisioned with a public cloud console

Explanation

A state file in Terraform does not represent a source of truth for resources provisioned with a public cloud console. The state file specifically tracks the resources provisioned using Terraform.

Your answer is correct

A state file represents a source of truth for resources provisioned with Terraform

Explanation

A state file in Terraform represents a source of truth for resources provisioned with Terraform. It stores the current state of the infrastructure and is used by Terraform to plan and apply changes to the infrastructure.

A state file represents the desired state expressed by the Terraform code files

Explanation

A state file in Terraform does not represent the desired state expressed by the Terraform code files. The state file stores the actual state of the infrastructure as managed by Terraform.

Overall explanation

The correct answer is: A state file represents a source of truth for resources provisioned with Terraform

Explanation:

The Terraform state file serves as the authoritative record of the infrastructure resources that Terraform manages. It contains detailed information about the current state of resources in the cloud or other infrastructure providers, as well as their mappings to the configuration in the Terraform code. This "source of truth" enables Terraform to:

- Track existing resources it manages.
- Detect and display changes during a `terraform plan`.
- Identify which resources to update, destroy, or leave untouched during a `terraform apply`.

Why the other options are incorrect:

- **"A state file can be used to schedule recurring infrastructure tasks"**: Terraform state is not used for task scheduling; it is for tracking resources and managing infrastructure.
- **"A state file represents a source of truth for resources provisioned with a public cloud console"**: The state file only tracks resources created or managed by Terraform, not those provisioned manually in the cloud console.
- **"A state file represents the desired state expressed by the Terraform code files"**: The state file represents the **current state** of resources, not the desired state (which is represented by the Terraform configuration files).

Summary:

The state file is a critical part of Terraform's operation and serves as the authoritative record of the resources it manages.

Question 6 **Incorrect**

Terraform configuration (including any module references) can contain only one Terraform provider type.

Correct answer

False

Explanation

This statement is correct. Terraform configuration files can include multiple provider blocks, each representing a different provider type. This flexibility allows users to manage resources from different cloud providers or services within a single Terraform configuration.

Your answer is incorrect

True

Explanation

This statement is incorrect. Terraform configuration files can contain multiple provider blocks, each specifying a different provider type. This allows users to manage resources from various cloud providers or services within the same configuration.

Overall explanation

Correct Answer: False

Explanation:

Terraform configurations can include multiple provider types. Each provider allows Terraform to interact with a specific API, enabling the management of resources across various platforms (e.g., AWS, Azure, Google Cloud, Kubernetes, etc.).

For example, you might have a configuration that manages resources in both AWS and Azure simultaneously by declaring both providers:

```
provider "aws" {
  region = "us-west-1"
}

provider "azurerm" {
  features = {}
}
```

Why True is incorrect:

The claim that a Terraform configuration can contain only one provider type is false. Terraform supports multi-cloud and multi-provider setups by allowing multiple provider blocks in a single configuration or referencing multiple modules with different providers.

Resources

[Terraform multi cloud provisioning](#)

Question 7Correct

Multiple team members are collaborating on infrastructure using Terraform and want to format their Terraform code following standard Terraform-style convention. How could they automatically ensure the code satisfies conventions?

Your answer is correct

Run the terraform fmt command during the code linting phase of your CI/CD process

Explanation

Running the `terraform fmt` command during the code linting phase of the CI/CD process automatically formats the Terraform code to adhere to standard Terraform-style conventions. This command ensures consistent formatting, such as indentation, spacing, and alignment of equal sign "=" characters, across all Terraform files, making it easier for team members to collaborate and maintain the codebase.

Run the terraform validate command prior to executing terraform plan or terraform apply

Explanation

Running the `terraform validate` command prior to executing `terraform plan` or `terraform apply` checks the syntax and configuration of the Terraform code for errors and warnings. While this command helps ensure the code is valid and can be executed, it does not specifically address formatting conventions or style guidelines. It is important to separate code validation from code formatting to maintain clean and readable Terraform code.

Manually apply two spaces indentation and align equal sign "=" characters in every Terraform file (*.tf)

Explanation

Manually applying two spaces indentation and aligning equal sign "=" characters in every Terraform file may not be practical or efficient, especially when multiple team members are collaborating on the infrastructure code. This approach is error-prone and time-consuming, as it relies on individual team members to consistently follow the formatting conventions, leading to inconsistencies and potential issues in the codebase.

Overall explanation

Correct Answer: Run the terraform fmt command during the code linting phase of your CI/CD process

Explanation:

The `terraform fmt` command automatically formats Terraform configuration files (*.tf) to align with Terraform's standard style conventions. It ensures consistent formatting, such as proper indentation, alignment of = characters, and whitespace management. Integrating this command into the CI/CD process ensures that all code contributed by team members follows a standard format, promoting readability and reducing stylistic conflicts.

Why the other options are incorrect:

- **Manually apply two spaces indentation and align equal sign "=" characters in every Terraform file (*.tf):** While you could manually format the code, it is time-consuming, error-prone, and unnecessary when `terraform fmt` exists to automate this process.
- **Run the terraform validate command prior to executing terraform plan or terraform apply:** The `terraform validate` command checks for syntax errors and ensures the configuration is syntactically valid, but it does not handle or enforce formatting. It focuses only on validating the correctness of the configuration.

Resources

[terraform fmt](#)

Question 8Correct

Define the purpose of state in Terraform.

State is used to store variables and quickly reuse existing code

Explanation

Storing variables and reusing existing code is not the primary purpose of state in Terraform. State is focused on mapping real-world resources to your configuration and tracking metadata to manage infrastructure.

Your answer is correct

State is used to map real world resources to your configuration and keep track of metadata

Explanation

State in Terraform is used to map the real-world resources that are being managed by your configuration. It keeps track of metadata such as resource IDs, attributes, and dependencies to ensure that your infrastructure stays in the desired state.

State is a method of codifying the dependencies of related resources

Explanation

While dependencies between resources are indeed managed in Terraform state, this choice does not fully capture the primary purpose of state, which is to map real-world resources to your configuration and track metadata.

State is used to enforce resource configurations that relate to compliance policies

Explanation

State in Terraform does not enforce resource configurations related to compliance policies. It is primarily used to track the state of resources and manage infrastructure as code.

Overall explanation

The correct answer is: State is used to map real-world resources to your configuration and keep track of metadata.

Explanation:

The **Terraform state** serves as a record of the resources managed by Terraform and is critical for its operation. It performs the following key functions:

- **Mapping Real-World Resources:** State maps the resources defined in your Terraform configuration to the actual resources created in your infrastructure (e.g., VMs, databases, networks).
- **Tracking Metadata:** It stores metadata such as resource dependencies, attributes, and the current state of resources, which helps Terraform understand the current infrastructure setup.
- **Enabling Plan Operations:** State allows Terraform to generate plans efficiently by comparing the current infrastructure (as recorded in the state) to the desired configuration defined in the code.
- **Ensuring Consistency:** It keeps the infrastructure in sync with the Terraform configuration, ensuring that changes made in the configuration are applied accurately.

Why the other options are incorrect:

- **"State is a method of codifying the dependencies of related resources":** While state tracks resource dependencies, codifying dependencies is achieved using Terraform's configuration, not its state file.
- **"State is used to enforce resource configurations that relate to compliance policies":** Compliance enforcement is handled through policy tools like Sentinel or manual validation, not Terraform state.
- **"State is used to store variables and quickly reuse existing code":** Variables are defined separately, and state is not intended for storing variables or reusing code.

Conclusion:

Terraform state plays a vital role in managing infrastructure and ensuring Terraform can track, plan, and apply changes accurately.

Resources

[Terraform language State](#)

Question 9 **Correct**

While attempting to deploy resources into your cloud provider using Terraform, you begin to see some odd behavior and experience slow responses. In order to troubleshoot you decide to turn on Terraform debugging. Which environment variables must be configured to make Terraform's logging more verbose?

TF_VAR_log_level

Explanation

TF_VAR_log_level is used to set the log level for variables in Terraform, not to configure the verbosity of Terraform logging.

TF_VAR_log_path

Explanation

TF_VAR_log_path is used to set the value of a variable in Terraform, not to configure the verbosity of Terraform logging.

Your answer is correct

TF_LOG

Explanation

TF_LOG is the correct environment variable that needs to be configured to make Terraform's logging more verbose. Setting TF_LOG to a value of TRACE will enable detailed logging for troubleshooting purposes.

TF_LOG_PATH

Explanation

TF_LOG_PATH is used to specify the file path where Terraform logs will be written. It is not used to control the verbosity of Terraform logging.

Overall explanation

Correct Answer:

TF_LOG

Explanation:

The **TF_LOG** environment variable is used to enable detailed logging in Terraform for debugging purposes. By setting **TF_LOG** to a specific verbosity level, such as **TRACE**, **DEBUG**, **INFO**, **WARN**, or **ERROR**, you can control the granularity of the logs Terraform outputs. These logs can help identify issues, such as API response errors or unexpected behavior during resource provisioning.

Example:

To enable debugging with detailed logs, you would configure the environment variable as follows:

```
bash
Copy codeexport TF_LOG=DEBUG
```

For the most detailed logs (including internal Terraform operations):

```
bash
Copy codeexport TF_LOG=TRACE
```

Why the Other Options Are Incorrect:

- **TF_LOG_PATH**: This environment variable is used to specify a file path to store the logs generated by **TF_LOG**, but it does not itself control the verbosity level.
- **TF_VAR_log_level**: This is not a recognized Terraform environment variable. Variables prefixed with **TF_VAR_** are used for input variables, not for logging configuration.
- **TF_VAR_log_path**: Similar to **TF_VAR_log_level**, this is not a valid Terraform environment variable. It would not configure Terraform's logging or verbosity.

The correct environment variable for enabling and controlling Terraform's logging verbosity is **TF_LOG**.

Resources

[Debugging Terraform](#)

Question 10Correct

When using Terraform to deploy resources into Azure, which scenarios are true regarding state files? (Choose two.)

Changing resources via the Azure Cloud Console records the change in the current state file

Explanation

Changing resources via the Azure Cloud Console does not automatically update the current state file. Terraform needs to be run with a plan or apply command to detect and reflect those changes in the state file.

Your selection is correct

Changing resources via the Azure Cloud Console does not update current state file

Explanation

Changing resources via the Azure Cloud Console does not automatically update the current state file. Terraform requires explicit commands like plan or apply to detect and update the state file based on the changes made outside of Terraform.

Your selection is correct

When you change a Terraform-managed resource via the Azure Cloud Console, Terraform updates the state file to reflect the change during the next plan or apply

Explanation

When you change a Terraform-managed resource via the Azure Cloud Console, Terraform updates the state file to reflect the change during the next plan or apply. This ensures that the state file remains up-to-date with the actual infrastructure configuration.

When you change a resource via the Azure Cloud Console, Terraform records the changes in a new state file

Explanation

When you change a resource via the Azure Cloud Console, Terraform does not create a new state file for each change. Instead, it updates the existing state file to reflect the modifications made to the infrastructure.

Overall explanation

Correct Answer:

- **When you change a Terraform-managed resource via the Azure Cloud Console, Terraform updates the state file to reflect the change during the next plan or apply.**
- **Changing resources via the Azure Cloud Console does not update the current state file.**

Explanation:

When you manually change resources via the Azure Cloud Console, Terraform does not immediately reflect those changes in its state file. However, during the next `terraform plan` or `terraform apply`, Terraform will detect the difference between the actual state in Azure and its state file. Terraform then updates the state file to reflect these changes, if needed.

- The first scenario is correct because Terraform will detect the changes during the next plan or apply and adjust the state file accordingly.
- The second scenario is incorrect because Terraform does not automatically update the state file for manual changes made through the Azure console. The state file needs to be refreshed or updated explicitly using `terraform refresh` or other methods.

Question 11 Correct

When using a remote backend or Terraform Cloud integration, where does Terraform save resource state?

In memory

Explanation

Terraform does not store resource state in memory when using a remote backend or Terraform Cloud integration. Storing state in memory would not allow for persistence and sharing of state information among team members.

In an environment variable

Explanation

Terraform does not save resource state in an environment variable when using a remote backend or Terraform Cloud integration. Environment variables are typically used for configuration settings and not for storing the state of infrastructure resources.

On the disk

Explanation

Terraform does not save resource state on the disk when using a remote backend or Terraform Cloud integration. The state is stored remotely to ensure consistency and collaboration among team members.

Your answer is correct

In the remote backend or Terraform Cloud

Explanation

When using a remote backend or Terraform Cloud integration, Terraform saves resource state in the remote backend or Terraform Cloud. This allows for centralized storage of state information, enabling collaboration, consistency, and version control among team members working on the same infrastructure.

Overall explanation

Correct Answer: In the remote backend or Terraform Cloud

Explanation:

When using a remote backend or integrating with Terraform Cloud, the resource state is securely stored and managed in the remote backend or Terraform Cloud. This allows for centralized state management, state locking, and collaboration across team members, ensuring that the state file is not stored locally on any individual machine.

Why the other options are incorrect:

- **On the disk:** This is true only for the default `local` backend. When using a remote backend or Terraform Cloud, the state file is not saved on the local disk.
- **In memory:** Terraform does load the state into memory during operations like `plan` or `apply`, but it is not persistently saved there. It is stored in the configured backend for future use.
- **In an environment variable:** Environment variables are used to provide secrets or configuration values but are not used to store the state. State storage is managed by the backend.

Resources

[Terraform Backend](#)

Question 12Incorrect

You want to share Terraform state with your team, store it securely, and provide state locking.

How would you do this? (Choose three.)

Using the s3 terraform backend. The dynamodb_field option is not needed.

Explanation

Using the s3 Terraform backend without the `dynamodb_field` option configured does not provide state locking capabilities. The `dynamodb_field` option is required to enable state locking when using an s3 backend for storing Terraform state.

Using the local backend.

Explanation

Using the local backend is not recommended for sharing Terraform state with a team as it does not provide the necessary security, state locking, and collaboration features that a remote backend like Terraform Cloud or Terraform Enterprise offers.

Correct selection

Using the consul Terraform backend.

Explanation

Using the Consul Terraform backend is another option for sharing Terraform state with your team. Consul provides a distributed key-value store that can be used to store Terraform state securely and provide state locking to prevent concurrent modifications.

Your selection is correct

Using an s3 terraform backend with an appropriate IAM policy and dynamodb_field option configured.

Explanation

Using an s3 Terraform backend with an appropriate IAM policy and `dynamodb_field` option configured allows you to securely store Terraform state in an s3 bucket and provide state locking using DynamoDB. This setup ensures that your team can collaborate on infrastructure changes without risking conflicts.

Your selection is correct

Using the remote Terraform backend with Terraform Cloud / Terraform Enterprise.

Explanation

Using the remote Terraform backend with Terraform Cloud / Terraform Enterprise allows you to securely store and share Terraform state with your team. It provides state locking to prevent conflicts when multiple team members are working on the same infrastructure.

Overall explanation

Correct Answers:

- Using the remote Terraform backend with Terraform Cloud / Terraform Enterprise.
- Using an s3 terraform backend with an appropriate IAM policy and `dynamodb_field` option configured.
- Using the consul Terraform backend.

Explanation:

Why these are correct:

- **Using the remote Terraform backend with Terraform Cloud / Terraform Enterprise:** Terraform Cloud and Terraform Enterprise offer a secure and centralized location for storing Terraform state files. These backends provide state locking, versioning, and collaboration features, making it easy for teams to work together safely. They also provide encrypted storage and allow multiple users to manage infrastructure without the risk of conflicting changes.
- **Using an s3 terraform backend with an appropriate IAM policy and dynamodb_field option configured:** When using Amazon S3 as a backend, state files are stored securely in an S3 bucket. The **DynamoDB** table is used for state locking, ensuring that only one user can make changes to the state at any given time. Configuring the appropriate IAM policy helps secure access to both the S3 bucket and DynamoDB table.
- **Using the consul Terraform backend:** Consul can also be used as a backend for Terraform, offering state locking and secure state storage, especially in environments where Consul is already in use for service discovery or distributed data storage.

Why others are incorrect:

- **Using the local backend:** The **local backend** stores Terraform state on the local filesystem, and it doesn't provide state sharing, state locking, or security features for teams working together. It's suitable for single-user scenarios but not for collaborative environments.
- **Using the s3 terraform backend. The dynamodb_field option is not needed:** This is incorrect because without configuring **DynamoDB** for state locking, multiple users could attempt to modify the state file at the same time, leading to potential state corruption or conflicts. The DynamoDB table is essential for enabling state locking in the S3 backend configuration.

Resources

[Terraform state backend](#)

Question 13 Correct

Sentinel policy-as-code is available in Terraform Enterprise.

False

Explanation

False. This statement is incorrect as Sentinel policy-as-code is a feature available in Terraform Enterprise. It provides a powerful tool for organizations to implement and enforce policies across their infrastructure deployments using code.

Your answer is correct

True

Explanation

True. Sentinel policy-as-code is indeed available in Terraform Enterprise. Sentinel allows users to define and enforce policies for their infrastructure as code to ensure compliance, security, and best practices are followed during the provisioning and management of resources.

Overall explanation

Correct Answer: True

Explanation:

Sentinel is a policy-as-code framework that is integrated with Terraform Enterprise (and Terraform Cloud) to enforce governance rules. It allows teams to write policies that can control and restrict certain actions, such as preventing deployments that don't meet security or compliance standards. This feature is available specifically in Terraform Enterprise, providing enhanced control over infrastructure deployments.

- Option "False" is incorrect because Sentinel is indeed available in Terraform Enterprise.

Resources

[Sentinel Docs](#)

Question 14 Incorrect

Which statement describes a goal of infrastructure as code?

A pipeline process to test and deliver software

Explanation

A pipeline process to test and deliver software is related to continuous integration and continuous delivery (CI/CD) practices, which focus on automating the software delivery process. While infrastructure as code can be part of a CI/CD pipeline, its primary goal is not specifically focused on testing and delivering software.

Your answer is incorrect

Write once, run anywhere

Explanation

"Write once, run anywhere" is a concept related to platform independence, where code can be written once and run on different platforms without modification. While this is a valuable goal in software development, infrastructure as code is more about managing infrastructure resources through code to ensure consistency and repeatability.

Correct answer

The programmatic configuration of resources

Explanation

The programmatic configuration of resources is a key goal of infrastructure as code. By defining infrastructure resources in code, users can automate the provisioning, configuration, and management of infrastructure components, leading to increased efficiency, consistency, and scalability in managing infrastructure. This approach allows for version control, collaboration, and repeatability in managing infrastructure resources.

Defining a vendor-agnostic API

Explanation

Defining a vendor-agnostic API refers to creating an interface that is not tied to any specific vendor or technology. While this can be a goal in software development, infrastructure as code is more about managing and provisioning infrastructure resources using code, rather than defining APIs.

Overall explanation

Correct Answer: The programmatic configuration of resources

Explanation:

The main goal of infrastructure as code (IaC) is to define and manage infrastructure through code, enabling you to programmatically configure and provision resources, making it easier to maintain, scale, and automate the deployment of infrastructure. This helps ensure consistency and repeatability in managing environments.

The other statements are more relevant to other areas of software development or general computing:

- "A pipeline process to test and deliver software" refers to continuous integration and continuous deployment (CI/CD) pipelines.
- "Defining a vendor-agnostic API" is about creating APIs that work across multiple platforms and vendors, not specifically about infrastructure management.
- "Write once, run anywhere" refers to creating software or code that can run on any platform without modification, typically in the context of software development rather than infrastructure.

Question 15 **Correct**

Which method for sharing Terraform configurations keeps them confidential within your organization, supports Terraform's semantic version constraints, and provides a browsable directory?

Generic git repository

Explanation

Using a generic git repository to share Terraform configurations may not keep them confidential within your organization, as git repositories can be accessed by anyone with the appropriate permissions. It also may not support Terraform's semantic version constraints and does not provide a browsable directory for easy navigation and management of modules.

Subfolder within a workspace

Explanation

Creating a subfolder within a workspace is not an ideal method for sharing Terraform configurations confidentially within your organization, as workspaces may not provide the necessary access controls to keep configurations secure. It also may not support Terraform's semantic version constraints and does not offer a browsable directory for easy management of modules.

Your answer is correct

Terraform Cloud/Terraform Enterprise private module registry

Explanation

Terraform Cloud/Terraform Enterprise private module registry is the correct choice because it allows organizations to keep their Terraform configurations confidential by restricting access to authorized users only. It also supports Terraform's semantic version

constraints, ensuring that modules are used correctly, and provides a browsable directory for easy navigation and management of modules within the organization.

Public Terraform Module Registry

Explanation

Public Terraform Module Registry is not the best choice for keeping Terraform configurations confidential within your organization, as modules in a public registry are accessible to anyone. While it does provide a browsable directory, it may not support Terraform's semantic version constraints for internal use cases.

Overall explanation

The correct answer is: Terraform Cloud/Terraform Enterprise private module registry.

Explanation:

- **Terraform Cloud/Terraform Enterprise private module registry** is designed to share Terraform configurations securely within an organization. It supports:
 - **Confidentiality:** Modules remain private to the organization.
 - **Semantic version constraints:** Ensures compatibility by allowing you to specify acceptable module versions using version constraints.
 - **Browsable directory:** Provides a user-friendly interface to browse available modules, making it easier for teams to find and reuse configurations.

Why Not the Others:

- **Generic git repository:** While it can store Terraform configurations, it lacks built-in support for semantic version constraints or a browsable directory tailored for Terraform.
- **Public Terraform Module Registry:** This is meant for sharing modules publicly, not for keeping configurations confidential within an organization.
- **Subfolder within a workspace:** This method lacks organization-wide visibility, versioning, and directory browsing features.

Conclusion:

Using the **Terraform Cloud/Terraform Enterprise private module registry** is the most secure, efficient, and feature-rich method for sharing Terraform configurations within an organization.

Resources

[Terraform Cloud Private module registry](#)

Question 16 **Incorrect**

You have decided to create a new Terraform workspace to deploy a development environment.

What is different about this workspace?

Correct answer

It has its own state file

Explanation

This choice is correct because each Terraform workspace has its own state file, which allows for isolation and separation of resources and configurations. This ensures that changes made in one workspace do not affect resources in another workspace.

It uses a different backend

Explanation

This choice is incorrect because the backend configuration is not specific to a workspace. The backend configuration determines where the Terraform state is stored, but each workspace can use the same backend configuration if needed. The state file is what distinguishes workspaces from each other.

It uses a different branch of code

Explanation

This choice is incorrect because workspaces do not necessarily use different branches of code. Workspaces are primarily used for managing state and isolating resources, while code branches are typically managed through version control systems like Git.

Your answer is incorrect

It pulls in a different terraform.tfvars file

Explanation

This choice is incorrect because the `terraform.tfvars` file is used to set variables for Terraform configurations and is not specific to a workspace. The state file, not the `tfvars` file, is what distinguishes workspaces from each other.

Overall explanation

Correct Answer: It has its own state file

Explanation: Each workspace in Terraform operates with a separate state file, allowing you to manage different environments or configurations independently. For example, you could have a workspace for development, another for production, and each would have its own state file, which keeps track of resources independently.

- "It pulls in a different `terraform.tfvars` file" is incorrect because workspaces do not automatically pull in a different `terraform.tfvars` file. The `terraform.tfvars` file can be defined separately, but it isn't specific to workspaces.
- "It uses a different branch of code" is incorrect as workspaces are not tied to specific branches in version control. Workspaces are isolated environments but use the same configuration code unless specifically modified.
- "It uses a different backend" is also incorrect because the backend remains the same across workspaces unless explicitly configured to be different. A backend is a remote location where Terraform stores state, and it isn't changed by switching workspaces.

Resources

[Workspaces](#)

Question 17 Correct

Which of these are secure options for storing secrets for connecting to a Terraform remote backend? **(Choose two.)**

Inside the backend block within the Terraform configuration

Explanation

Storing secrets inside the backend block within the Terraform configuration is not a secure option as the configuration files are often version-controlled and can be accessed by anyone with access to the repository. This can lead to potential security risks and exposure of sensitive information.

A variable file

Explanation

Storing secrets in a variable file is not a secure option for connecting to a Terraform remote backend. Variable files are often included in version-controlled repositories and can be accessed by unauthorized users, leading to potential security vulnerabilities.

Your selection is correct

Defined in a connection configuration outside of Terraform

Explanation

Defining secrets in a connection configuration outside of Terraform is a secure option for storing secrets for connecting to a Terraform remote backend. By keeping the secrets outside of the Terraform configuration files, you can better control access to the sensitive information and reduce the risk of exposure.

Your selection is correct

Defined in Environment variables

Explanation

Defining secrets in environment variables is a secure option for storing secrets for connecting to a Terraform remote backend. Environment variables are not typically stored in version-controlled files and provide a more secure way to manage sensitive information.

Overall explanation

Correct Answers: Defined in Environment variables and Defined in a connection configuration outside of Terraform

Explanation:

- **Defined in Environment variables:** Storing secrets in environment variables is a secure and recommended way to manage sensitive information. Terraform supports using environment variables (e.g., `TF_VAR_name` for variable injection) to securely pass secrets at runtime, without hardcoding them into the configuration files.
- **Defined in a connection configuration outside of Terraform:** Using external tools or configuration management systems (e.g., AWS Secrets Manager, HashiCorp Vault) to manage and inject secrets into the Terraform runtime is another secure option. This keeps sensitive information outside the Terraform codebase and state files.
- **Inside the backend block within the Terraform configuration:** This is not secure because secrets defined in the backend block will be included in the Terraform configuration files, which might be shared or exposed.

- **A variable file:** While possible, this is less secure unless the file is encrypted or stored in a protected location. Plaintext variable files can easily be exposed if the repository or storage location is compromised.

Question 18 **Incorrect**

When should you run `terraform init` command?

After you run terraform apply for the first time in a new Terraform project and before you run terraform plan

Explanation

Running terraform init after you run terraform apply for the first time in a new Terraform project is not the correct timing. Terraform init should be executed before running terraform plan to initialize the working directory and download any necessary plugins.

Correct answer

After you start coding a new Terraform project and before you run terraform plan for the first time

Explanation

Running terraform init after you start coding a new Terraform project and before you run terraform plan for the first time is the correct timing. Terraform init initializes the working directory and downloads any necessary plugins to ensure that the Terraform configuration is ready for planning and applying.

After you run terraform plan for the first time in a new Terraform project and before you run terraform apply

Explanation

Running terraform init after you run terraform plan for the first time in a new Terraform project is not the correct timing. Terraform init should be executed before running terraform plan to initialize the working directory and download any necessary plugins.

Your answer is incorrect

Before you start coding a new Terraform project

Explanation

Running terraform init before you start coding a new Terraform project is not necessary. Terraform init should be executed after you start coding the project and before running terraform plan to set up the environment and prepare for the execution of Terraform commands.

Overall explanation

Correct Answer: After you start coding a new Terraform project and before you run terraform plan for the first time

Explanation:

The `terraform init` command is used to initialize a Terraform working directory. It sets up the backend, downloads the necessary provider plugins, and prepares the project for execution. It must be run:

- **After coding a new Terraform project:** This ensures Terraform can properly configure the backend and fetch any required providers defined in the configuration files before you proceed with running `terraform plan` or `terraform apply`.
- **Before running terraform plan for the first time:** Without initialization, Terraform cannot build an execution plan as it lacks the necessary providers and backend setup.

Why the others are incorrect:

- **After you run terraform apply and before terraform plan:** Initialization is required before the plan stage, not after applying changes.
- **After terraform plan and before terraform apply:** If you haven't initialized the project, Terraform won't even be able to run the plan.
- **Before you start coding a new Terraform project:** You need configuration files present in the directory for `terraform init` to have any meaningful effect. Running it before coding has no utility.

Resources

[Terraform init](#)

Question 19 **Correct**

You are working on some new application features and you want to spin up a copy of your production deployment to perform some quick tests. In order to avoid having to configure a new state backend, what open source Terraform feature would allow you create multiple states but still be associated with your current code?

Terraform modules

Explanation

Terraform modules are used to organize and reuse Terraform configurations. While modules can help with code organization and reusability, they do not directly address the need for managing multiple states or state backends.

Terraform data sources

Explanation

Terraform data sources are used to fetch data from an external source during the Terraform execution. They are not related to managing multiple states or state backends in Terraform.

Terraform local values

Explanation

Terraform local values are used to define variables within a Terraform configuration file. They are not related to managing multiple states or state backends in Terraform.

None of the above

Explanation

None of the above choices directly address the specific requirement of managing multiple states while still being associated with your current code in Terraform.

Your answer is correct

Terraform workspaces

Explanation

Terraform workspaces allow you to manage multiple states within the same configuration. Each workspace can have its own state file, allowing you to work with different environments or configurations without the need to configure separate state backends. This makes it the correct choice for creating multiple states associated with your current code.

Overall explanation

Correct Answer: Terraform workspaces

Explanation:

Terraform workspaces allow you to manage multiple environments or versions of your infrastructure using the same code base. Each workspace has its own state file, so you can create isolated copies of your infrastructure, like spinning up a copy of your production deployment for testing purposes. This allows you to perform tests without needing to configure a new backend or change your code.

- **Terraform data sources** are used to query information from existing resources but do not manage multiple states or environments.
- **Terraform local values** are used to assign local variables within the configuration and do not manage separate state files or environments.
- **Terraform modules** are reusable components of Terraform code, but they do not handle managing different states or environments by themselves.

Workspaces are designed specifically to handle multiple environments with isolated state files while maintaining a single set of configuration files.

Resources

[Workspaces](#)

Question 20Incorrect

You are creating a reusable Terraform configuration and want to include a `billing_dept` tag so your Finance team can track team-specific spending on resources. Which of the following `billing_dept` variable declarations will allow you to do this?

```
variable "billing_dept" {
```

```
    optional = true
```

```
}
```

Explanation

This declaration is incorrect because the `"optional"` argument is not a valid type for a variable declaration in Terraform. The correct syntax for defining a variable as optional is to set the `"default"` argument to null or an empty string.

Your answer is incorrect


```
variable "billing_dept" {
```

```
    type = default
```

```
}
```

Explanation

This declaration is incorrect because the "type" argument should specify the data type of the variable, not "default." Additionally, using "default" as a type is not a valid syntax for defining a variable in Terraform.

Correct answer

```
variable "billing_dept" {
```

```
    default = ""
```

```
}
```

Explanation

This declaration is correct because it defines the "billing_dept" variable with a default value of an empty string. This allows the Finance team to track team-specific spending on resources by including this variable in the Terraform configuration.

```
variable "billing_dept" {
```

```
    type = optional(string)
```

```
}
```

Explanation

This declaration is incorrect because the "type" argument should specify the data type of the variable, such as "string" in this case. Using "optional(string)" is not a valid syntax for defining a variable in Terraform.

Overall explanation

Correct Answer: `variable "billing_dept" { default = "" }`

Explanation:

Declaring a variable with a default value of an empty string (`default = ""`) ensures that the `billing_dept` variable is optional. If no value is explicitly provided, Terraform will use the default value. This approach makes the configuration flexible and reusable while avoiding errors due to missing variable values.

- `variable "billing_dept" { optional = true }`: This syntax is invalid in Terraform, as `optional` is not a supported attribute for variable declarations.
- `variable "billing_dept" { type = optional(string) }`: This is also invalid because `optional(string)` is not a valid type in Terraform. Type constraints must use standard types such as `string`, `list`, or `map`.
- `variable "billing_dept" { type = default }`: This is invalid syntax, as `default` is not a valid type in Terraform variable declarations.

Using `default = ""` is the correct method to make the variable optional and allow it to be omitted during execution without causing errors.

Resources

[Terraform Input variables](#)

Question 21 Correct

Which of these commands makes your code more human readable?

```
terraform validate
```

Explanation

The `terraform validate` command is used to validate the syntax of the Terraform configuration files and check for any errors. While it ensures that the code is valid, it does not directly impact the readability of the code for humans.

```
terraform plan
```

Explanation

The `terraform plan` command is used to create an execution plan that shows what Terraform will do when you apply the configuration. While it is essential for understanding the changes that will be made to the infrastructure, it does not directly enhance the human readability of the code.

terraform output

Explanation

The `terraform output` command is used to extract the values of output variables defined in the Terraform configuration. It helps in viewing the output of the infrastructure provisioning process but does not specifically improve the human readability of the code.

Your answer is correct

terraform fmt

Explanation

The `terraform fmt` command is used to automatically format the Terraform configuration files according to a standard style guide. It helps in standardizing the code layout, indentation, and formatting, making it more consistent and easier for humans to read and understand.

Overall explanation

Correct Answer: terraform fmt

Explanation: The `terraform fmt` command automatically formats your Terraform configuration files to adhere to standard style conventions, making the code more consistent and human-readable. It ensures proper indentation, spacing, and alignment, making it easier for others to read and understand the code.

- `terraform validate` is used to check if your Terraform configuration is syntactically valid but does not affect the readability of the code itself.
- `terraform output` is used to display the values of outputs defined in the configuration but does not impact the formatting of the code.
- `terraform plan` shows what changes Terraform will make to the infrastructure but does not format the code or enhance its readability.

Resources

[terraform fmt command](#)

Question 22Correct

When should you write Terraform configuration files for existing infrastructure that you want to start managing with Terraform?

Terraform will generate the corresponding configuration files for you

Explanation

Terraform does not automatically generate corresponding configuration files for existing infrastructure that you want to manage. You need to write the configuration files yourself to define the infrastructure as code and leverage Terraform's capabilities for provisioning, updating, and managing the infrastructure.

Your answer is correct

Before you run terraform import

Explanation

Writing Terraform configuration files before running `terraform import` allows you to define the desired state of the existing infrastructure and map it to Terraform resources. This approach ensures that the infrastructure is managed consistently and in line with your configuration, avoiding any discrepancies between the actual state and the desired state.

After you run terraform import

Explanation

Writing Terraform configuration files after running `terraform import` may lead to inconsistencies between the actual state of the infrastructure and the desired state defined in the configuration files. It is best practice to define the configuration before importing the infrastructure to ensure accurate management and alignment with your infrastructure requirements.

You can import infrastructure without corresponding Terraform code

Explanation

While it is possible to import existing infrastructure without corresponding Terraform code, it is recommended to have Terraform configuration files in place to manage and maintain the infrastructure effectively. Without Terraform code, it may be challenging to track changes, apply updates, or collaborate with team members on infrastructure changes.

Overall explanation

Correct Answer: Before you run terraform import

Explanation:

Before using `terraform import` to bring existing infrastructure under Terraform management, you must manually write the corresponding Terraform configuration files. These files describe the infrastructure resources that already exist. Terraform does not automatically generate configuration files during the import process; it only adds the resource to the state file.

Here's why other answers are incorrect:

- **You can import infrastructure without corresponding Terraform code** is incorrect because Terraform requires a matching resource definition in the configuration files before importing it.
- **Terraform will generate the corresponding configuration files for you** is incorrect because Terraform does not create configuration files automatically. You are responsible for writing them.
- **After you run terraform import** is incorrect because the resource needs to exist in the configuration beforehand. If it doesn't, Terraform will fail to import the resource.

Resources

[Importing Infrastructure](#)

Question 23Correct

Which of the following is not a valid Terraform variable type?

Your answer is correct

array

Explanation

The 'array' variable type is not a valid Terraform variable type. Terraform does not have a specific 'array' variable type, as lists are used to represent sequences of values instead.

list

Explanation

The 'list' variable type in Terraform is used to represent a sequence of values. It is a valid variable type that allows you to define a list of elements within a Terraform configuration.

string

Explanation

The 'string' variable type in Terraform is used to represent a sequence of characters. It is a valid variable type that allows you to define string values within a Terraform configuration.

map

Explanation

The 'map' variable type in Terraform is used to represent a collection of key-value pairs. It is a valid variable type that allows you to define a mapping of keys to values within a Terraform configuration.

Overall explanation

The correct answer is: array

Explanation:

In Terraform, the valid variable types are:

1. **list**: Represents an ordered collection of values, typically all of the same type.
2. **map**: Represents a collection of key-value pairs.
3. **string**: Represents a single textual value.

Terraform does **not** have a type called **array**. Instead, Terraform uses **list** for what is commonly referred to as an array in other programming contexts.

Example of valid types:

```
variable "example_list" {  
  type = list(string)
```

```

    }

    variable "example_map" {
      type = map(string)
    }

    variable "example_string" {
      type = string
    }

```

Summary:

- **list**, **map**, and **string** are valid Terraform types.
- **array** is not a valid type in Terraform.

Resources

[Expression Types and Values](#)

Question 24Correct

What are some benefits of using Sentinel with Terraform Cloud/Terraform Enterprise? (**Choose three.**)

Sentinel Policies can be written in HashiCorp Configuration Language (HCL)

Explanation

Sentinel policies are written in the Sentinel policy language, not HCL. While HCL is used for Terraform configurations, Sentinel has its own syntax and constructs.

You can check out and check in cloud access keys

Explanation

Checking out and checking in cloud access keys is not a benefit of using Sentinel with Terraform Cloud/Terraform Enterprise.

Sentinel focuses on policy enforcement and governance aspects, rather than managing access keys or credentials. It is important to handle access keys securely and separately from policy enforcement mechanisms.

Your selection is correct

Policy-as-code can enforce security best practices

Explanation

Using Sentinel with Terraform Cloud/Terraform Enterprise allows you to enforce security best practices by implementing policy-as-code. This means you can define and enforce rules and guidelines for infrastructure provisioning, ensuring that security measures are consistently applied across all deployments.

Your selection is correct

You can restrict specific configurations on resources like "CIDR=0.0.0.0/0" not allowed

Explanation

With Sentinel, you can restrict specific configurations on resources, such as disallowing configurations like "CIDR=0.0.0.0/0". This helps in preventing misconfigurations or insecure settings that could potentially lead to security vulnerabilities or breaches.

Your selection is correct

You can enforce a list of approved AWS AMIs

Explanation

Sentinel enables you to enforce a list of approved AWS AMIs (Amazon Machine Images) that can be used in your infrastructure deployments. This helps in maintaining consistency, compliance, and security by ensuring that only authorized and vetted AMIs are utilized.

Overall explanation

Correct Answers:

Policy-as-code can enforce security best practices

You can restrict specific configurations on resources like "CIDR=0.0.0.0/0" not allowed

You can enforce a list of approved AWS AMIs

Explanation:

Sentinel is a policy-as-code framework integrated with Terraform Cloud and Terraform Enterprise, allowing organizations to enforce governance, compliance, and security policies directly within their workflows.

Why these are correct:

- **Policy-as-code can enforce security best practices:** Sentinel enables organizations to define and enforce security, compliance, and operational policies within Terraform workflows, ensuring best practices are adhered to.
- **You can restrict specific configurations on resources like "CIDR=0.0.0.0/0" not allowed:** Sentinel can validate Terraform plans and enforce rules such as disallowing overly permissive network configurations (e.g., public access through `CIDR=0.0.0.0/0`).
- **You can enforce a list of approved AWS AMIs:** Sentinel policies can specify allowed values, such as a predefined list of approved Amazon Machine Images (AMIs), ensuring consistency and compliance.

Why others are incorrect:

- **Sentinel Policies can be written in HashiCorp Configuration Language (HCL):** Sentinel policies are written in the Sentinel policy language, not HCL. While HCL is used for Terraform configurations, Sentinel has its own syntax and constructs.
- **You can check out and check in cloud access keys:** This is unrelated to Sentinel. Managing cloud access keys is typically handled through secret management tools or cloud provider IAM systems, not Sentinel policies.

Resources

[Sentinel](#)

Question 25Incorrect

If you manually destroy infrastructure, what is the best practice reflecting this change in Terraform?

Run terraform import

Explanation

Running ``terraform import`` is used to bring existing resources under Terraform management, not to reflect manual destruction of infrastructure. It is not the appropriate action to take in this scenario.

Correct answer

Remove the resource definition from your file and run terraform apply -refresh-only

Explanation

Removing the resource definition from your Terraform configuration file and running ``terraform apply -refresh-only`` is the correct approach when manually destroying infrastructure. This ensures that Terraform recognizes the change and updates its state accordingly without affecting other resources.

Manually update the state file

Explanation

Manually updating the state file is not the recommended practice when destroying infrastructure in Terraform. The state file should accurately reflect the actual state of the infrastructure managed by Terraform, so manual updates can lead to inconsistencies and potential issues.

Your answer is incorrect

It will happen automatically

Explanation

Infrastructure destruction does not happen automatically in Terraform. It is important to follow the correct procedures to ensure that Terraform's state file accurately reflects the actual state of the infrastructure.

Overall explanation

Correct Answer: Remove the resource definition from your file and run terraform apply -refresh-only

Explanation:

If you manually destroy infrastructure (such as deleting a resource directly through the cloud provider's console), Terraform will no longer track this resource in its state file, causing a mismatch between the actual infrastructure and the state file.

The best practice to reflect this change in Terraform is to:

1. **Remove the resource definition** from your Terraform configuration.
2. Run `terraform apply -refresh-only` to refresh the state file and remove the missing resource from Terraform's state.

This ensures Terraform's state is updated without attempting to recreate or manage resources that no longer exist.

Why the other options are not correct:

- **Manually update the state file** – Directly editing the state file is not recommended as it could lead to inconsistency or corruption in your state. It's safer to let Terraform manage the state.
- **Run terraform import** – `terraform import` is used when you want to import existing infrastructure into Terraform's management. If you've manually deleted infrastructure, `terraform import` would not be the correct approach, as you don't want Terraform to manage the resource anymore.
- **It will happen automatically** – Terraform does not automatically detect when resources are manually deleted outside of Terraform, so it won't automatically adjust the state file or configuration. You need to run the appropriate commands to reflect these changes in the state file.

Conclusion:

After manually destroying infrastructure, the correct approach is to remove the corresponding resource from the Terraform configuration file and run `terraform apply -refresh-only` to update the state file. This ensures your infrastructure and state remain consistent.

Resources

[Terraform refresh command](#)

Question 26 **Incorrect**

Both Terraform Cloud and Terraform Enterprise support policy as code (Sentinel).

Your answer is incorrect

False

Explanation

False. This statement is incorrect as both Terraform Cloud and Terraform Enterprise do support policy as code through Sentinel. Sentinel provides a powerful framework for implementing and enforcing policies across the infrastructure deployment process in a consistent and automated manner.

Correct answer

True

Explanation

True. Both Terraform Cloud and Terraform Enterprise support policy as code through Sentinel. Sentinel allows users to define and enforce policies to govern infrastructure as code deployments, ensuring compliance, security, and operational best practices are followed.

Overall explanation

The correct answer is: True

Explanation:

Both **Terraform Cloud** and **Terraform Enterprise** support **policy as code** using **Sentinel**, HashiCorp's embedded policy-as-code framework. Sentinel allows organizations to enforce fine-grained, automated policies on Terraform runs to ensure compliance, security, and operational standards before resources are provisioned.

- In **Terraform Cloud**, Sentinel is available in the **Team & Governance** and **Business** tiers.
- In **Terraform Enterprise**, Sentinel is also included, providing organizations with the ability to enforce policies across their infrastructure provisioning workflows.

Resources

[Sentinel Terraform](#)

Question 27 **Incorrect**

A developer accidentally launched a VM (virtual machine) outside of the Terraform workflow and ended up with two servers with the same name. They don't know which VM Terraform manages but do have a list of all active VM IDs.

Which of the following methods could you use to discover which instance Terraform manages?

Correct answer

Run terraform state list to find the names of all VMs, then run terraform state show for each of them to find which VM ID Terraform manages

Explanation

Running terraform state list to find the names of all VMs and then running terraform state show for each of them is the correct method to discover which instance Terraform manages. By listing the state of all VMs and then inspecting each one individually, you can identify which VM Terraform is currently managing based on the VM ID.

Use terraform refresh/code to find out which IDs are already part of state

Explanation

Using terraform refresh/code to find out which IDs are already part of the state is not the most efficient method to discover which instance Terraform manages. While terraform refresh updates the state file with the real-world infrastructure, it may not directly help in identifying the specific instance managed by Terraform.

Run terraform taint/code on all the VMs to recreate them

Explanation

Running terraform taint/code on all the VMs to recreate them is not the correct method to discover which instance Terraform manages. Tainting resources will mark them for recreation, but it does not provide information on which VM instance Terraform is currently managing.

Your answer is incorrect

Update the code to include outputs for the ID of all VMs, then run terraform plan to view the outputs

Explanation

Updating the code to include outputs for the ID of all VMs and then running terraform plan to view the outputs may provide information on the VM IDs, but it does not directly help in identifying which instance Terraform manages. Outputs are used to display information after the execution of Terraform commands, not to determine the managed instances.

Overall explanation

Correct Answer:

Run `terraform state list` to find the names of all VMs, then run `terraform state show` for each of them to find which VM ID Terraform manages.

Explanation:

When you have a list of active VM IDs and need to determine which one is managed by Terraform, you can use the Terraform state file, which maps real-world infrastructure to your Terraform configuration. Here's how to approach this:

1. **Use `terraform state list`:**
This command lists all the resources currently being tracked in Terraform's state. From this, you can identify the Terraform resource names associated with VMs.
2. **Use `terraform state show`:**
For each VM listed in the Terraform state, use `terraform state show <resource_name>` to view detailed metadata about that resource, including its VM ID. Compare the ID against your list to find the match.

This method ensures you accurately identify which VM is managed by Terraform without modifying or recreating resources.

Why Other Methods Are Incorrect:

- **Run `terraform taint` on all the VMs to recreate them:**
This is risky and unnecessary. Tainting marks resources for destruction and recreation, which could disrupt existing infrastructure.
- **Update the code to include outputs for the ID of all VMs, then run `terraform plan`:**
While adding outputs can help track VM IDs, this approach would not immediately show existing resource IDs already managed by Terraform unless you apply the updated code and refresh the state.
- **Use `terraform refresh` to find out which IDs are already part of state:**
While `terraform refresh` updates the state to match the real-world infrastructure, it doesn't directly show which VM ID is managed. You'd still need to inspect the state with `terraform state show` to get the details.

This approach avoids unnecessary changes and focuses on using Terraform's built-in state inspection tools effectively.

Resources

[terraform state list](#)

Question 28Incorrect

Please fill the blank field(s) in the statement with the right words.

Some backends like s3 support ____.

Your answer is incorrect

versioning

Correct answer

state locking

Explanation

Correct Answer: Some

Explanation:

Not all backends in Terraform support state locking. **Some backends**, such as **Amazon S3 (with DynamoDB for locking)**, **Terraform Cloud**, and **Consul**, provide support for state locking to prevent multiple users from making conflicting changes to the same state file. State locking ensures that only one process can modify the state at a time, avoiding potential corruption or conflicts.

Resources

[State Storage and Locking](#)

Question 29Correct

The Terraform binary version and provider versions must match each other in a single configuration.

Your answer is correct

False

True

Overall explanation

Correct Answer:

False

Explanation:

The Terraform binary version and provider versions **do not need to match** each other. Terraform is designed to manage providers independently through version constraints. This allows you to use different provider versions regardless of the specific Terraform binary version, as long as they are compatible.

Key Details:

- Provider Versioning:**
 - Providers are defined in the configuration file with version constraints (e.g., `>= 3.0.0`, `< 4.0.0`).
 - Terraform uses the provider version specified in the configuration or the version recorded in the `.terraform.lock.hcl` file.
- Terraform Binary Compatibility:**
 - The Terraform binary enforces compatibility with provider versions but does not require them to match directly.
 - You can upgrade the Terraform binary without upgrading the providers or vice versa.
- Provider Independence:**
 - Providers are downloaded separately from the Terraform Registry and stored in the `.terraform` directory.
 - This independence ensures flexibility and modularity in managing infrastructure.

Example:

You might use Terraform 1.5.0 with the AWS provider version 5.0.0, as long as the provider is compatible with the Terraform binary. The `.terraform.lock.hcl` file ensures consistent provider versions across team members and environments.

Why "True" Is Incorrect:

Requiring the Terraform binary and provider versions to match would limit flexibility and make version management unnecessarily complex. Terraform was specifically designed to avoid this constraint.

Question 30 **Correct**

Which of these actions will prevent two Terraform runs from changing the same state file at the same time?

Refresh the state after running Terraform

Explanation

Refreshing the state after running Terraform does not prevent two Terraform runs from changing the same state file at the same time. State refreshing is used to update the state file with the current state of the infrastructure, but it does not provide locking mechanisms to prevent concurrent state changes.

Your answer is correct

Configure state locking for your state backend

Explanation

Configuring state locking for your state backend is the correct choice because it allows you to prevent two Terraform runs from changing the same state file at the same time. State locking ensures that only one Terraform run can write to the state file at any given time, preventing conflicts and data corruption.

Delete the state before running Terraform

Explanation

Deleting the state before running Terraform will not prevent two Terraform runs from changing the same state file at the same time. Deleting the state will remove the existing state file, but it does not implement any locking mechanisms to ensure exclusive access to the state file during Terraform runs.

Run Terraform with parallelism set to 1

Explanation

Running Terraform with parallelism set to 1 may reduce the likelihood of concurrent state changes, but it does not provide a robust solution to prevent two Terraform runs from changing the same state file at the same time. State locking is the recommended approach to ensure exclusive access to the state file during Terraform runs.

Overall explanation

Correct Answer:

Configure state locking for your state backend

Explanation:

State locking is a critical feature that prevents concurrent modifications to the state file, ensuring that two Terraform runs do not accidentally make conflicting changes. When state locking is enabled, Terraform will "lock" the state file before making changes and "unlock" it after the operation is complete. This feature is especially important in collaborative environments where multiple users or automation processes might operate on the same Terraform state.

- **State locking is supported by remote backends** such as S3 (with DynamoDB for locking), Terraform Cloud, Consul, and others. These backends ensure that only one operation can modify the state file at a time.

Why the Other Options Are Incorrect:

- **"Refresh the state after running Terraform"**: Refreshing the state simply updates Terraform's understanding of the current state of resources but does not prevent concurrent state modifications.
- **"Delete the state before running Terraform"**: Deleting the state file would not prevent concurrent modifications; it would instead cause Terraform to lose track of the infrastructure, leading to potential configuration drift or errors.
- **"Run Terraform with parallelism set to 1"**: Setting parallelism to 1 reduces the number of concurrent resource operations within a single run, but it does not prevent multiple runs from accessing and modifying the state file simultaneously.

State locking is the only mechanism specifically designed to address this issue effectively.

Resources

[State Locking](#)

Question 31 **Incorrect**

When using multiple configurations of the same Terraform provider, what meta-argument must be included in any non-default provider configurations?

Correct answer

alias

Explanation

The `alias` meta-argument is necessary for distinguishing between multiple configurations of the same Terraform provider. By providing a unique alias for each configuration, Terraform can differentiate between them and apply the correct configuration to the corresponding resources.

name

Explanation

The `name` meta-argument is used to assign a specific name to a resource within Terraform, but it is not a mandatory meta-argument for handling multiple configurations of the same provider. The `name` attribute is more focused on resource naming conventions rather than provider configuration differentiation.

id

Explanation

The `id` meta-argument is used to uniquely identify resources within Terraform, but it is not specifically required for managing multiple configurations of the same provider. The `id` is typically generated by Terraform and used for resource tracking and referencing.

Your answer is incorrect

depends_on

Explanation

The `depends_on` meta-argument is used to define explicit dependencies between resources in Terraform, ensuring that certain resources are created or destroyed in a specific order. It is not required for specifying multiple configurations of the same provider.

Overall explanation

Correct Answer:

alias

Explanation:

When using multiple configurations of the same Terraform provider in a single configuration, the **alias** meta-argument is required for any non-default provider configurations. This allows you to reference and differentiate between multiple instances of the same provider.

Details:

1. Default Provider Configuration:

By default, a provider block is used as-is without any special designation:

```
provider "aws" {
  region = "us-east-1"
}
```

2. Additional Provider Configurations:

If you want to use multiple configurations of the same provider (e.g., managing resources in multiple AWS regions), you need to specify an alias for the additional configurations:

```
provider "aws" {
  region = "us-east-1"
}

provider "aws" {
  alias   = "west"
  region = "us-west-1"
}
```

3. Using the Aliased Provider:

To reference the non-default provider configuration, the **provider** argument in the resource block must specify the alias:

```
resource "aws_instance" "example" {
  provider = aws.west
}
```

```
ami      = "ami-12345678"
instance_type = "t2.micro"
}
```

4. **Why `alias` Is Required:**

The `alias` uniquely identifies the non-default provider configurations. Without it, Terraform would not be able to differentiate between multiple configurations of the same provider.

Why Other Options Are Incorrect:

- **`depends_on:`**
This is used to specify dependencies between resources and is unrelated to provider configurations.
- **`id:`**
The `id` attribute is used to identify specific resources or outputs, not provider configurations.
- **`name:`**
There is no `name` meta-argument for providers in Terraform. The `alias` meta-argument serves this purpose.

Resources

[alias: Multiple Provider Configurations](#)

Question 32 **Incorrect**

Infrastructure as Code (IaC) can be stored in a version control system along with application code.

Your answer is incorrect

False

Explanation

False. Storing Infrastructure as Code (IaC) separately from application code can lead to inconsistencies, difficulties in tracking changes, and challenges in maintaining version control. It is recommended to store IaC alongside application code in a version control system for better management and coordination.

Correct answer

True

Explanation

True. Storing Infrastructure as Code (IaC) in a version control system along with application code is a best practice in DevOps. This allows for versioning, tracking changes, collaboration, and ensuring consistency between infrastructure and application code.

Overall explanation

Correct Answer: True

Explanation: Infrastructure as Code (IaC) can indeed be stored in a version control system (VCS) alongside application code. This practice allows teams to version, track, and manage infrastructure changes in the same way they handle application code. By storing IaC in a VCS like Git, you ensure that both the application and the infrastructure code can be managed collaboratively, with the ability to roll back changes, review history, and coordinate between development and operations teams.

- Storing IaC in a VCS provides many benefits like collaboration, change tracking, and maintaining consistent configurations across environments. This is considered a best practice in modern DevOps workflows.

Question 33 **Correct**

You have a list of numbers that represents the number of free CPU cores on each virtual cluster:

numcpus = [18, 3, 7, 11, 2]

What Terraform function could you use to select the largest number from the list?

top(numcpus)

Explanation

There is no built-in Terraform function called `top()` that can be used to select the largest number from a list. This choice is incorrect as it does not align with the functionality required in the question.

ceil(numcpus)

Explanation

The `ceil()` function in Terraform is used to round a number up to the nearest whole number. It is not suitable for selecting the largest number from a list of numbers like in this scenario.

`high[numcpus]`

Explanation

The syntax `high[numcpus]` is not a valid Terraform function or expression. It does not exist in Terraform and cannot be used to select the largest number from a list of numbers.

Your answer is correct

`max(numcpus)`

Explanation

The correct Terraform function to select the largest number from a list is `max()`. In this case, using `max(numcpus)` will return the largest number from the `numcpus` list, which is 18.

Overall explanation

The correct answer is: `max(numcpus)`

Explanation:

- The `max()` function in Terraform is used to return the largest value from a list of numbers or a set of arguments.
- In this case, using `max(numcpus)` will evaluate the list `[18, 3, 7, 11, 2]` and return 18, the largest number.

Why the other options are incorrect:

- `ceil(numcpus)`: The `ceil()` function in Terraform applies to a single floating-point number, rounding it up to the nearest whole number. It does not work with lists.
- `top(numcpus)`: This is not a valid Terraform function.
- `high[numcpus]`: This is also not a valid Terraform function or syntax.

Correct Usage:

```
numcpus = [18, 3, 7, 11, 2]
largest = max(numcpus)
```

This would set `largest` to 18.

Resources

[max Function](#)

Question 34Correct

Which backend does the Terraform CLI use by default?

Terraform Cloud

Explanation

Terraform Cloud is a remote backend option that can be used with the Terraform CLI, but it is not the default backend. Terraform Cloud provides additional features such as remote state management, collaboration tools, and run history, but it is not the default backend for Terraform operations.

API

Explanation

The Terraform CLI does not use an API as the default backend. APIs are used for communication between different software systems and services, but they are not the default backend for Terraform operations.

Your answer is correct

Local

Explanation

The Terraform CLI uses a local backend by default. This means that the state file is stored on the local file system of the user running Terraform. The local backend is simple and easy to set up, making it the default choice for most Terraform users.

Remote

Explanation

The Terraform CLI does not use a remote backend as the default. Remote backends are used for storing state files in a centralized location to enable collaboration and version control, but they are not the default backend for Terraform operations.

HTTP

Explanation

The Terraform CLI does not use HTTP as the default backend. HTTP is a protocol used for communication over the internet, but it is not the default backend for Terraform operations.

Overall explanation

The correct answer is: Local.

Explanation:

By default, the **Terraform CLI** uses the **local backend**. The local backend stores the Terraform state file (`terraform.tfstate`) locally on the machine where Terraform is executed. This is the simplest backend and does not require additional configuration.

Key Features of the Local Backend:

- **Local Storage:** The state file is stored as a file on the local filesystem.
- **No Remote Features:** It does not support remote state sharing or locking unless explicitly configured with additional tools (e.g., a locking mechanism like Consul or DynamoDB).
- **Default Behavior:** Unless a specific backend is explicitly defined in the Terraform configuration, the local backend is used.

Why Other Options Are Incorrect:

- **API:** This is not a backend type in Terraform.
- **Remote:** This is a generic term that refers to backends that store state remotely (e.g., S3, Terraform Cloud), but it is not the default.
- **Terraform Cloud:** While Terraform Cloud can be used as a remote backend, it requires explicit configuration in the `terraform` block.
- **HTTP:** This is a supported backend for state storage, but it is not the default.

Conclusion:

The **local backend** is the default backend for Terraform CLI, providing basic state storage locally without any additional setup.

Resources

[Local backend](#)

Question 35Correct

You created infrastructure outside of the Terraform workflow that you now want to manage using Terraform. Which command brings the infrastructure into Terraform state?

`terraform init`

Explanation

The `terraform init` command is used to initialize a Terraform working directory, but it does not bring existing infrastructure into Terraform state. It is typically used at the beginning of a new Terraform project to set up the environment.

Your answer is correct

`terraform import`

Explanation

The `terraform import` command is used to bring existing infrastructure under Terraform management by importing it into the Terraform state. This command allows you to manage resources that were created outside of Terraform using its configuration files.

`terraform refresh`

Explanation

The `terraform refresh` command is used to update the state file with the real-world infrastructure. It does not bring existing infrastructure into Terraform state. It is used to reconcile the state Terraform knows about with the real-world infrastructure.

`terraform get`

Explanation

The `terraform get` command is used to download and update modules mentioned in the root module. It does not bring existing infrastructure into Terraform state. It is used to fetch modules from the Terraform registry or a version control system.

Overall explanation

Correct Answer: terraform import

Explanation:

When you create infrastructure outside of the Terraform workflow and want to bring that infrastructure under Terraform management, the correct approach is to use the `terraform import` command. This command allows you to import existing resources into Terraform's state so that Terraform can begin managing them.

For example, if you created a virtual machine outside of Terraform and now want to manage it, you would run `terraform import` followed by the resource type and its ID to import it into Terraform.

Why the other options are not correct:

- **terraform init** – This command initializes your working directory for Terraform, downloading provider plugins, and setting up the backend. It does not import existing resources into the state file.
- **terraform get** – This command is used to download and update modules in your configuration, not for importing existing infrastructure into Terraform's state.
- **terraform refresh** – This command updates the Terraform state by querying the infrastructure to get the current state of managed resources. It does not import resources that Terraform is not yet managing.

Conclusion:

To bring existing infrastructure into Terraform's management, you use the `terraform import` command.

Resources

[Terraform cli import](#)

Question 36Incorrect

You want to define multiple data disks as nested blocks inside the resource block for a virtual machine.

What Terraform feature would help you define the blocks using the values in a variable?

Count arguments

Explanation

Count arguments in Terraform are used to create multiple instances of a resource or module based on a specified count value. While count arguments can help with creating multiple resources, they are not ideal for defining nested blocks inside a resource block using variable values.

Your answer is incorrect

Local values

Explanation

Local values in Terraform are used to assign a value to a variable within a module or configuration. They are not directly related to defining nested blocks inside a resource block using variable values.

Collection functions

Explanation

Collection functions in Terraform are used to manipulate lists, maps, and sets of values. While they can help with data manipulation, they are not specifically designed to define nested blocks inside a resource block using variable values.

Correct answer

Dynamic blocks

Explanation

Dynamic blocks in Terraform allow you to create multiple nested blocks inside a resource block based on a list or map variable. This feature is specifically designed to dynamically generate nested blocks using variable values, making it the correct choice for defining multiple data disks for a virtual machine.

Overall explanation

The correct answer is: Dynamic blocks

Explanation:

Dynamic blocks in Terraform allow you to programmatically generate nested blocks within a resource by iterating over values, such as a list or map provided by a variable. This feature is particularly useful for defining multiple nested blocks, such as data disks for a virtual machine, without duplicating code for each block manually.

For example:

```
resource "virtual_machine" "example" {
  name = "vm-example"

  dynamic "data_disk" {
    for_each = var.disks
    content {
      size      = data_disk.value.size
      name      = data_disk.value.name
      type      = data_disk.value.type
    }
  }
}
```

Here:

- The `dynamic` block generates multiple `data_disk` blocks.
- `for_each` iterates over the values in the `var.disks` variable.

Why others are incorrect:

- **Local values:** Local values are used to simplify expressions but do not generate nested blocks.
- **Collection functions:** These are used for transformations or filtering on lists or maps but do not create nested blocks.
- **Count arguments:** `count` can duplicate a resource or block, but it does not handle complex nested block generation efficiently. For nested blocks, `dynamic` is better suited.

Resources

[Dynamic Blocks in Terraform](#)

Question 37Correct

How would you refer to the indexing instance from the below configuration?

```
resource "aws_instance" "web" {
  ...
  for_each = {
    "terraform": "value1",
    "resource": "value2",
    "indexing": "value3",
    "example": "valued",
  }
}
```

`aws_instance.web.indexing`

Explanation

The correct way to refer to the indexing instance in the given configuration is by using dot notation: `aws_instance.web.indexing`.

However, when using the `for_each` meta-argument, the key should be accessed using square brackets and quotes, not dot notation.

`aws_instance-web["indexing"]`

Explanation

The syntax `aws_instance-web["indexing"]` is incorrect in Terraform. The correct way to refer to the instance with the key "indexing" in this configuration is by using `aws_instance.web["indexing"]` with square brackets and quotes.

Your answer is correct

`aws_instance.web["indexing"]`

Explanation

The correct way to refer to the indexing instance in the provided configuration is `aws_instance.web["indexing"]`. This syntax correctly accesses the specific instance with the key "indexing" using the `for_each` meta-argument in Terraform.

`aws_instance["web"]["indexing"]`

Explanation

The correct syntax for referring to an instance using the `for_each` meta-argument in Terraform is `aws_instance.web["indexing"]`. Using square brackets and the key name within quotes is the appropriate way to access the specific instance in this configuration.

Overall explanation

Correct Answer: `aws_instance.web["indexing"]`

Explanation:

In Terraform, when you use `for_each` with a map, each instance of the resource is created with a key from the map. To access a specific instance, you reference it using the syntax `<resource_type>.<resource_name>["<key>"]`.

In this case:

- The resource type is `aws_instance`.
- The resource name is `web`.
- The key in the `for_each` map is `"indexing"`.

To access the instance created with the key `"indexing"`, you write `aws_instance.web["indexing"]`.

Why others are incorrect:

The syntax `aws_instance["web"] ["indexing"]` is not valid because Terraform does not allow you to reference the resource name (`web`) inside brackets.

The syntax `aws_instance.web.indexing` is incorrect because keys in a `for_each` map must be accessed using brackets, not dot notation. Dot notation is only used to access resource attributes.

The syntax `aws_instance-web["indexing"]` is invalid as it does not follow Terraform's syntax rules. The use of `|` and the structure is not supported in Terraform configurations.

When working with a `for_each` map, the proper way to reference an instance is `aws_instance.web["indexing"]`.

Resources

[Referring to Instances](#)

Question 38Correct

The public Module Registry is free to use.

False

Explanation

False. The public Module Registry is actually free to use. Users can leverage the modules available in the registry to accelerate their Terraform development without incurring any charges.

Your answer is correct

True

Explanation

True. The public Module Registry provided by HashiCorp is indeed free to use. Users can access a wide range of pre-built modules for Terraform without any cost, making it a valuable resource for infrastructure as code development.

Overall explanation

The correct answer is: True

Explanation:

The **public Module Registry** provided by HashiCorp is free to use. It allows users to discover, share, and use pre-built Terraform modules for a variety of infrastructure components and providers. These modules can help speed up infrastructure provisioning by providing reusable configurations.

While the public registry is free, it is important to note that private registries (for internal organizational use) might require additional setup or licensing, depending on your Terraform implementation (e.g., Terraform Cloud or Enterprise).

Resources

[Terraform Registry Publishing](#)

Question 39Correct

Using the `terraform state rm` command against a resource will destroy it.

True

Explanation

True. Using the `terraform state rm` command against a resource will remove it from the Terraform state file, but it will not actually destroy the resource itself. The resource will still exist in the cloud provider unless explicitly destroyed through Terraform apply or other means.

Your answer is correct

False

Explanation

False. Using the `terraform state rm` command against a resource will only remove it from the Terraform state file, but it will not destroy the resource itself. The resource will still exist in the cloud provider unless explicitly destroyed through Terraform apply or other means.

Overall explanation

The correct answer is: False.

Explanation:

The `terraform state rm` command **does not destroy the resource itself**. Instead, it **removes the specified resource from the Terraform state file**. This means Terraform will no longer manage that resource, but the resource itself will remain intact in the infrastructure.

Key Details:

- **Purpose of `terraform state rm`:** It is used to disconnect a resource from Terraform's state file without affecting the actual resource.
- **Impact on Infrastructure:** The resource will continue to exist in the cloud provider or infrastructure, but Terraform will no longer track or manage it.
- **Common Use Case:** This command is often used when a resource needs to be managed manually or through another process, or if there's an issue with the state file that requires cleanup.

What Actually Destroys a Resource:

To destroy a resource, you would use the `terraform destroy` command or update your configuration to remove the resource, then run `terraform apply`.

Conclusion:

The `terraform state rm` command only removes the resource from the state file and does not delete it from the infrastructure.

Resources

[terraform state rm](#)

Question 40Correct

Select the command that doesn't cause Terraform to refresh its state.

`terraform plan`

Explanation

The `terraform plan` command is used to generate an execution plan showing what Terraform will do when you call `terraform apply`. This command will cause Terraform to refresh its state by comparing the current state with the desired state and displaying the planned changes.

Your answer is correct

terraform state list

Explanation

The `terraform state list` command is used to list all resources in the Terraform state. This command does not cause Terraform to refresh its state but simply displays the existing resources in the state file.

terraform apply

Explanation

The `terraform apply` command is used to apply the changes required to reach the desired state of the configuration. This command will cause Terraform to refresh its state by comparing the current state with the desired state and making any necessary changes.

terraform destroy

Explanation

The `terraform destroy` command is used to destroy all the resources defined in the Terraform configuration. This command will cause Terraform to refresh its state by removing all resources and updating the state file accordingly.

Overall explanation

Correct Answer: terraform state list

Explanation:

The `terraform state list` command is used to list the resources tracked in the current state file, but it does not trigger a refresh of the state. It simply displays the current state as it is.

- `terraform apply`, `terraform destroy`, and `terraform plan` all involve refreshing the state, as Terraform needs to compare the actual state with the desired configuration to determine the changes to apply or destroy.

Question 41 Correct

You must initialize your working directory before running `terraform validate`.

Your answer is correct

True

Explanation

True. It is necessary to initialize the working directory with `terraform init` before running `terraform validate`. This command ensures that all necessary plugins and modules are downloaded and configured correctly before validating the Terraform configuration.

False

Explanation

False. This statement is incorrect. Initializing the working directory with `terraform init` is a prerequisite for running `terraform validate` to ensure that the Terraform configuration is valid and can be applied successfully.

Overall explanation

Correct Answer: True

Explanation:

You **must initialize your working directory before running** `terraform validate` because `terraform validate` checks the syntax and configuration of your Terraform files. Initializing your working directory with `terraform init` ensures that the necessary provider plugins, modules, and configurations are downloaded and set up, allowing `terraform validate` to function correctly.

If the directory is not initialized, Terraform will not be able to validate the configuration properly since it hasn't set up the necessary files or dependencies required for the operation.

Why This is True:

- `terraform init` initializes the working directory by downloading the necessary providers and modules, setting up the backend configuration, and preparing the environment for Terraform to work with.

- **terraform validate** will check for issues in the syntax and structure of the Terraform files, but it relies on the initialization step to ensure that all required dependencies (like providers and modules) are available and can be used during validation.

Conclusion:

Always run `terraform init` before `terraform validate` to ensure that Terraform is properly set up and ready to validate your configuration correctly.

Resources

[Terraform cli commands validate](#)

Question 42Correct

You must use different Terraform commands depending on the cloud provider you use.

True

Explanation

This statement is incorrect. Terraform commands are consistent across different cloud providers. The same commands such as ``terraform init``, ``terraform plan``, ``terraform apply``, and ``terraform destroy`` are used regardless of the cloud provider being used.

Your answer is correct

False

Explanation

This statement is correct. Terraform commands are designed to be cloud-agnostic, meaning they can be used with any cloud provider without the need for different commands. This consistency allows users to maintain a standardized workflow regardless of the cloud environment they are working with.

Overall explanation

Correct Answer:

False

Explanation:

Terraform is a cloud-agnostic tool, meaning its commands remain consistent regardless of the cloud provider (e.g., AWS, Azure, GCP, etc.) or the type of infrastructure you are managing. Commands such as `terraform init`, `terraform plan`, `terraform apply`, and `terraform destroy` are universal across all providers. The provider-specific details are managed through Terraform's provider plugins and configurations, not the commands themselves.

This allows users to interact with multiple cloud providers using the same set of Terraform commands, ensuring a consistent workflow across diverse environments.

Question 43Incorrect

You want to define a single input variable to capture configuration values for a server. The values must represent memory as a number, and the server name as a string.

Which variable type could you use for this input?

Correct answer

Object

Explanation

The Object variable type would be the correct choice for capturing configuration values for a server that include memory as a number and server name as a string. Object variables allow you to define a complex data structure with multiple attributes of different types, making it suitable for storing both a number and a string in a single input variable.

List

Explanation

Using a List variable type would not be suitable for capturing configuration values for a server that include memory as a number and server name as a string. List variables are used to represent a collection of values of the same type, such as a list of strings or numbers. In this case, you need a variable type that can hold different types of values.

Terraform does not support complex input variables of different types

Explanation

Terraform does support complex input variables of different types, such as using the Object variable type to capture configuration values for a server that include memory as a number and server name as a string. It is important to choose the appropriate variable type based on the data structure you need to represent in your Terraform configuration.

Your answer is incorrect

Map

Explanation

Using a Map variable type would not be the most appropriate choice for capturing configuration values for a server that include memory as a number and server name as a string. Map variables are used to represent key-value pairs, where both the key and value are of the same type. In this case, you need a variable type that can hold different types of values.

Overall explanation

Correct Answer: Object

Explanation:

An **object** is the most suitable variable type for this input because it allows you to define a complex variable with multiple attributes of different types. For example, you can use an object to capture the server configuration, specifying **memory** as a number and **server_name** as a string. Here's how you can define and use it in Terraform:

```
variable "server_config" {
  type = object({
    memory      = number
    server_name = string
  })
}

# Example usage:
resource "example_server" "main" {
  memory      = var.server_config.memory
  server_name = var.server_config.server_name
}
```

Why the other options are incorrect:

- **List:** A list can only store multiple items of the same type (e.g., a list of strings or a list of numbers). It cannot mix data types, making it unsuitable for this use case.
- **Map:** A map stores key-value pairs, but all values must be of the same type. This makes it less flexible for mixed-type configurations like **memory** (number) and **server_name** (string).
- **Terraform does not support complex input variables of different types:** This is incorrect because Terraform supports complex types like objects and tuples that allow mixed data types.

Resources

[Input Variables](#)

[Object Data type](#)

Question 44Correct

Which command should you run to check if all code in a Terraform configuration that references multiple modules is properly formatted without making changes?

terraform fmt -list -recursive

Explanation

The command "terraform fmt -list -recursive" is not a valid Terraform command for checking if code in a configuration that references multiple modules is properly formatted. The "-list" flag is not a valid option for the "terraform fmt" command.

terraform fmt -write=false

Explanation

The command "terraform fmt -write=false" is used to check if all code in a Terraform configuration is properly formatted without making any changes. However, this command does not specifically target configurations that reference multiple modules.

Your answer is correct

terraform fmt -check -recursive

Explanation

The command "terraform fmt -check -recursive" is the correct choice for checking if all code in a Terraform configuration that references multiple modules is properly formatted without making changes. The "-check" flag is used to perform a check without making any changes, and the "-recursive" flag ensures that all nested directories are included in the check.

terraform fmt -check

Explanation

The command "terraform fmt -check" is not the most suitable option for checking if all code in a Terraform configuration that references multiple modules is properly formatted without making changes. The "-check" flag is used to perform a check without making any changes, but the "-recursive" flag is necessary to include all nested directories when checking configurations with multiple modules.

Overall explanation

Correct Answer: terraform fmt -check -recursive

Explanation:

The `terraform fmt` command checks and formats Terraform configuration files to ensure consistent formatting. To verify if all code referencing multiple modules is properly formatted **without making changes**, you use the `-check` flag. Adding the `-recursive` flag ensures that the command checks all directories, including subdirectories containing module code.

- `terraform fmt -write=false` is incorrect because the `-write=false` flag does not exist in Terraform.
- `terraform fmt -list -recursive` is incorrect because the `-list` flag only lists files that would be changed but does not check formatting.
- `terraform fmt -check` is partially correct but does not recursively check formatting in subdirectories.

Resources

[terraform fmt command](#)

Question 45Correct

What features stops multiple users from operating on the Terraform state at the same time?

Remote backends

Explanation

Remote backends in Terraform are used to store the state file remotely, but they do not inherently prevent multiple users from operating on the state at the same time. State locking is a separate feature designed specifically for this purpose.

Your answer is correct

State locking

Explanation

State locking is the feature in Terraform that prevents multiple users from operating on the state at the same time. It ensures that only one user can make changes to the state file at any given time, preventing conflicts and data corruption.

Provider constraints

Explanation

Provider constraints are not directly related to preventing multiple users from operating on the Terraform state at the same time. Providers in Terraform are responsible for managing the lifecycle of a resource, not for state locking.

Version control

Explanation

Version control systems like Git are used to manage changes to code and configurations over time, but they do not provide the functionality to prevent multiple users from operating on the Terraform state at the same time. State locking is the specific feature in Terraform designed for this purpose.

Overall explanation

Correct Answer: State locking

Explanation:

State locking is a feature in Terraform that prevents multiple users or processes from modifying the same state file simultaneously. It ensures consistency and avoids conflicts or corruption of the state file when multiple operations (e.g., `terraform plan`, `terraform apply`) are performed concurrently.

- **Provider constraints** are unrelated; they control the version of providers used but don't manage state conflicts.
- **Remote backends** provide shared storage for the state file but don't inherently prevent simultaneous access. State locking is often implemented as part of remote backends like S3 with DynamoDB or Terraform Cloud.
- **Version control** helps manage configuration files but doesn't manage state or prevent simultaneous operations on the state file.

Resources

[Terraform state locking](#)

Question 46Correct

You are writing a child Terraform module which provisions an AWS instance. You want to make use of the IP address returned in the root configuration. You name the instance resource "main".

Which of these is the correct way to define the output value using HCL2?

```
output "instance_ip_addr" {
```

```
    value = "${aws_instance.main.private_ip}"
```

```
}
```

Explanation

This choice uses incorrect interpolation syntax by enclosing the reference to the private IP address of the AWS instance in "\${}". The correct way to reference attributes of resources in HCL2 is by using the dot notation as shown in choice A.

```
output "instance_ip_addr" {
```

```
    return aws_instance.main.private_ip
```

```
}
```

Explanation

This choice uses the incorrect keyword "return" instead of "value" to define the output value. In HCL2, the correct keyword to assign a value to an output is "value". Additionally, the syntax used to reference the private IP address of the AWS instance is incorrect.

```
output "aws_instance.instance_ip_addr" {
```

```
    value = "${main.private_ip}"
```

```
}
```

Explanation

This choice contains a syntax error with the use of the incorrect characters in the output name and the interpolation syntax. The correct way to reference the private IP address of the AWS instance named "main" is by using the dot notation as shown in choice A.

Your answer is correct

```
output "instance_ip_addr" {
```

```
    value = aws_instance.main.private_ip
```

```
}
```

Explanation

This choice correctly defines an output named "instance_ip_addr" with the value of the private IP address of the AWS instance named "main". The syntax used is valid for HCL2 and follows the correct format for referencing the attribute of the AWS instance resource.

Overall explanation

Correct Answer: `output "instance_ip_addr" { value = aws_instance.main.private_ip }`

Explanation:

The correct syntax for defining an output in HCL2 (HashiCorp Configuration Language) follows the structure:

- `output "name" { value = <expression> }`

Why this is correct:

- `output "instance_ip_addr" { value = aws_instance.main.private_ip }` correctly defines an output block with the name `instance_ip_addr` and assigns the private IP of the `aws_instance.main` resource to the value. This is valid HCL2 syntax.

Why the others are incorrect:

- `output "aws_instance.instance_ip_addr" { value = "${main.private_ip}" }:`
 - Incorrect because the output name includes a `.` which is not allowed.
 - `${...}` interpolation is also unnecessary in HCL2 for simple expressions.
- `output "instance_ip_addr" { value = "${aws_instance.main.private_ip}" }:`
 - While this would have worked in earlier versions of Terraform, `${...}` string interpolation is no longer required for straightforward expressions in HCL2. You can directly use `aws_instance.main.private_ip`.
- `output "instance_ip_addr" { return aws_instance.main.private_ip }:`
 - Incorrect because `return` is not a valid keyword in HCL2 for outputs. Instead, `value` must be used.

By using `value = aws_instance.main.private_ip`, the private IP of the instance resource `main` is correctly referenced and exposed as an output.

Resources

[Output Values](#)

Question 47Incorrect

Which feature is not included in Terraform Cloud's free tier?

Correct answer

Audit logging

Explanation

Audit logging is not included in Terraform Cloud's free tier. Audit logging allows users to track and monitor changes made to their infrastructure, providing visibility into who made changes and when.

Workspace

Explanation

Workspaces are included in Terraform Cloud's free tier. They allow users to organize and manage their infrastructure code in separate environments.

Private module registry

Explanation

Private module registry is included in Terraform Cloud's free tier. It allows users to store and share Terraform modules privately within their organization.

Your answer is incorrect

Remote state management

Explanation

Remote state management is included in Terraform Cloud's free tier. It enables users to store and access their Terraform state files remotely, providing a centralized location for state management.

Overall explanation

Correct Answer: Audit logging

Explanation:

Audit logging is not included in Terraform Cloud's free tier. This feature is available only in higher tiers (such as business or premium plans) and is designed to track detailed logs of user actions and system activities for compliance and security purposes.

Why others are included in the free tier:

- **Workspace:** Terraform Cloud's free tier allows unlimited workspaces to manage infrastructure for different projects or environments.
- **Remote state management:** The free tier provides remote state storage and management, enabling collaboration and preventing state file conflicts.
- **Private module registry:** The free tier includes a private module registry, allowing teams to share and reuse Terraform modules within their organization.

Audit logging, being a premium feature, is excluded from the free plan, as it targets enterprises requiring advanced monitoring and compliance capabilities.

Resources

[Terraform audit logs](#)

Question 48 **Correct**

It is best practice to store secret data in the same version control repository as your Terraform configuration.

True

Explanation

Storing secret data in the same version control repository as your Terraform configuration is not considered a best practice. Doing so can expose sensitive information to unauthorized users who have access to the repository. It is recommended to use secure secret management tools or services to store and manage secret data separately from the Terraform configuration.

Your answer is correct

False

Explanation

This choice is correct because storing secret data in the same version control repository as your Terraform configuration is a security risk. It is not recommended to store sensitive information like passwords, API keys, or other credentials in plain text within the repository. Instead, utilize secure secret management solutions to protect this data from unauthorized access.

Overall explanation

Correct Answer: False

Explanation:

It is not recommended to store secret data (such as passwords, API keys, or other sensitive information) in the same version control repository as your Terraform configuration. Storing secrets in version-controlled repositories could expose them to unauthorized access and compromise the security of your infrastructure.

Best practices for managing secrets include:

- Storing secrets in secure external services like AWS Secrets Manager, HashiCorp Vault, or Azure Key Vault.
- Using environment variables to inject sensitive data into Terraform configurations.
- Leveraging encryption and key management to safeguard sensitive information.

Hence, it is considered a security risk to store secret data directly within the version control repository that holds your Terraform configurations.

Question 49 **Correct**

Before you can use Terraform's remote backend, you must first execute terraform init.

Your answer is correct

True

Explanation

True. Before using Terraform's remote backend, you must initialize the configuration by running the "terraform init" command. This command downloads the necessary plugins and modules, sets up the backend configuration, and prepares the working directory for Terraform to use.

False

Explanation

False. This statement is incorrect. In order to utilize Terraform's remote backend, it is essential to first execute the "terraform init" command. This step is crucial for setting up the necessary configurations and preparing the environment for remote backend usage.

Overall explanation

Correct Answer: True

Explanation:

Before using Terraform's remote backend (or any backend), you must first run `terraform init`. This command initializes the working directory containing your Terraform configuration files and sets up the backend, downloading necessary provider plugins and configuring the remote backend if specified. Without this initialization, Terraform won't know where or how to store the state or configure the backend.

- Option "False" is incorrect because `terraform init` is a required step to initialize the backend, among other tasks like setting up provider plugins.

Resources

[Terraform Backend](#)

Question 50Incorrect

Which of these statements about Terraform Enterprise workspaces is false?

Your answer is incorrect

They have role-based access controls

Explanation

This statement is true. Terraform Enterprise workspaces have role-based access controls, which enable administrators to define and manage user permissions within the workspace environment.

Plans and applies can be triggered via version control system integrations

Explanation

This statement is true. Plans and applies in Terraform Enterprise workspaces can be triggered via version control system integrations, allowing for automated infrastructure changes based on code commits.

Correct answer

You must use the CLI to switch between workspaces

Explanation

This statement is false. In Terraform Enterprise, you do not need to use the CLI to switch between workspaces. Workspaces can be managed and switched within the Terraform Enterprise UI, making it easier to work with multiple environments.

They can securely store cloud credentials

Explanation

This statement is true. Terraform Enterprise workspaces can securely store cloud credentials, allowing for secure access to cloud resources during Terraform operations.

Overall explanation

The correct answer is: You must use the CLI to switch between workspaces

Explanation:

This statement is **false** because in Terraform Enterprise, workspaces are not limited to the CLI for switching. Workspaces in Terraform Enterprise can be managed and accessed through its web interface, where users can view, create, and switch between workspaces seamlessly without relying on the CLI.

Why the other options are true:

- **"They can securely store cloud credentials"**: Terraform Enterprise workspaces securely store sensitive data, including cloud provider credentials, using Vault-backed secrets management.
- **"Plans and applies can be triggered via version control system integrations"**: Terraform Enterprise integrates with version control systems (e.g., GitHub, GitLab, Bitbucket). When code changes are pushed to a linked repository, it can automatically trigger plan and apply operations.

- **"They have role-based access controls"**: Terraform Enterprise provides robust role-based access controls, allowing organizations to manage who can view, plan, and apply changes within specific workspaces.

Summary:

Workspaces in Terraform Enterprise are designed for flexibility and ease of use, including GUI-based interactions, making the CLI-only switching statement incorrect.

Question 51 **Incorrect**

You can develop a custom provider to manage its resources using Terraform.

Your answer is incorrect

False

Explanation

False. Developing a custom provider to manage resources using Terraform is a valid and supported practice. By creating a custom provider, users can define and manage resources that are specific to their infrastructure needs, providing flexibility and customization in resource management.

Correct answer

True

Explanation

True. In Terraform, custom providers can be developed to manage resources that are not supported by the built-in providers. This allows users to extend Terraform's capabilities and manage a wider range of resources within their infrastructure.

Overall explanation

Correct Answer:

True

Explanation:

Terraform supports the development of custom providers, allowing users to manage resources not covered by existing providers. This flexibility enables organizations to integrate Terraform with proprietary or less common platforms, systems, or APIs.

Details:

1. **Custom Provider Definition:**
A custom provider is a plugin that interacts with an API or system to provision and manage resources via Terraform.
2. **When to Develop a Custom Provider:**
 - No existing provider supports the platform or API you want to manage.
 - Your infrastructure includes proprietary or in-house systems requiring integration.
3. **How Custom Providers Work:**
 - Custom providers use Terraform's **Provider Plugin SDK** or **Terraform Plugin Framework** (Go programming language).
 - They define resource types, data sources, and logic for managing the system.
 - The provider is compiled into a binary and referenced in the Terraform configuration.
4. **Development Steps:**
 - Install the Terraform Plugin SDK.
 - Define the schema for resources and data sources.
 - Implement CRUD (Create, Read, Update, Delete) operations for resources.
 - Test and validate the provider.
5. **Use Case Examples:**
 - Managing internal company systems, such as a custom database or CI/CD tool.
 - Integrating with less common cloud platforms or on-premises tools.

Why "False" Is Incorrect:

It is entirely possible to develop custom providers with Terraform. HashiCorp provides tools, SDKs, and documentation to help developers build and use their own providers.

Resources

[Implement a provider with the Terraform Plugin Framework](#)

Question 52 **Incorrect**

You add a new resource to an existing Terraform configuration, but do not update the version constraint in the configuration. The existing and new resources use the same provider. The working directory contains a `.terraform-lock.hcl` file.

How will Terraform choose which version of the provider to use?

Terraform will use the latest version of the provider available at the time you provision your new resource

Explanation

This choice is incorrect because Terraform does not automatically use the latest version of the provider available at the time of provisioning new resources. Instead, it relies on the version recorded in the lock file to ensure consistency and predictability in the configuration.

Your answer is incorrect

Terraform will check your state file to determine the provider version to use

Explanation

This choice is incorrect because Terraform does not determine the provider version to use based on the state file. The state file stores the state of the resources and their configurations, but it does not dictate the provider version.

Correct answer

Terraform will use the version recorded in your lock file

Explanation

This choice is correct because Terraform will prioritize the version recorded in the lock file to maintain consistency in the provider version across all resources. This ensures that the configuration remains stable and predictable.

Terraform will use the latest version of the provider for the new resource and the version recorded in the lock file to manage existing resources

Explanation

This choice is incorrect because Terraform will use the version recorded in the lock file to manage both existing and new resources. The lock file ensures that the provider version remains consistent across all resources in the configuration.

Overall explanation

Correct Answer:

Terraform will use the version recorded in your lock file

Explanation:

Terraform uses the `.terraform-lock.hcl` file to ensure consistent provider versions across runs. This file records the exact version of each provider used in the configuration. When you add a new resource to an existing configuration and do not update the version constraint in the configuration, Terraform will continue to use the provider version recorded in the lock file for **all resources** (existing and new). This behavior ensures stability and reproducibility in your infrastructure management.

Why other options are incorrect:

- **"Terraform will use the latest version of the provider for the new resource and the version recorded in the lock file to manage existing resources"**: Terraform does not mix provider versions between resources in the same configuration. All resources will use the version specified in the lock file.
- **"Terraform will check your state file to determine the provider version to use"**: The state file records information about managed resources but does not dictate the provider version; this is controlled by the lock file.
- **"Terraform will use the latest version of the provider available at the time you provision your new resource"**: Terraform does not automatically update providers to the latest version unless explicitly instructed via a `terraform init -upgrade` or by modifying the version constraints in the configuration and updating the lock file.

Resources

[Dependency Installation Behavior](#)

Question 53 **Correct**

From which of these sources can Terraform import modules?

GitHub Repository

Explanation

Terraform can import modules directly from a GitHub repository. This is particularly useful when the module is hosted on GitHub and needs to be easily accessible and version-controlled.

Terraform Module Registry

Explanation

The Terraform Module Registry is an official repository of Terraform modules that can be easily searched, shared, and used by the Terraform community. Importing modules from the Terraform Module Registry ensures that the modules are vetted, maintained, and easily accessible.

Local path

Explanation

Terraform allows importing modules from a local path on the user's machine. This can be useful for testing or development purposes when the module is not yet published or shared with others.

Your answer is correct

All of the above

Explanation

All of the above choices are valid sources from which Terraform can import modules. Whether it's from a local path, a GitHub repository, or the Terraform Module Registry, Terraform provides flexibility in importing modules from various sources to meet different use cases and requirements.

Overall explanation

Correct Answer:

All of the above

Explanation:

Terraform can import modules from multiple sources, including:

- **Local path:** You can use a local path to a module that exists on your local machine. This allows you to reuse code that is stored on your local file system.

Example:

```
module "example" {
  source = "./path/to/module"
}
```

- **GitHub Repository:** Terraform allows you to import modules directly from a GitHub repository by specifying the GitHub URL as the source. This is useful when you want to use publicly available or private modules stored on GitHub.

Example:

```
module "example" {
  source = "github.com/username/repository//path/to/module"
}
```

- **Terraform Module Registry:** The **Terraform Module Registry** is the official source for reusable modules. Terraform provides access to thousands of community-contributed modules that you can use directly in your configurations.

Example:

```
module "example" {
  source = "terraform-google-modules/vpc/google"
}
```

Thus, Terraform can import modules from a **local path**, a **GitHub repository**, or the **Terraform Module Registry**, making it flexible for a wide range of use cases.

Resources

[Module Sources](#)

Question 54Correct

Which command adds existing resources into Terraform state?

terraform refresh

Explanation

The `terraform refresh` command is used to update the state file with the real-world infrastructure. It does not add existing resources into the Terraform state, but it updates the state with the current state of the resources.

Your answer is correct

terraform import

Explanation

The `terraform import` command is used to import existing infrastructure into Terraform. It adds existing resources into the Terraform state, allowing you to manage them using Terraform configurations.

terraform plan

Explanation

The `terraform plan` command is used to create an execution plan for Terraform. It shows what actions Terraform will take when you apply your configuration. It does not add existing resources into the Terraform state.

All of these

Explanation

While `terraform import` is the correct choice for adding existing resources into Terraform state, the other commands (`terraform init`, `terraform plan`, `terraform refresh`) do not perform the same function. Therefore, not all of the choices listed are correct for adding existing resources into Terraform state.

terraform init

Explanation

The `terraform init` command is used to initialize a Terraform working directory, download necessary plugins, and prepare the working directory for other commands. It does not add existing resources into the Terraform state.

Overall explanation

Correct Answer: terraform import

Explanation:

The `terraform import` command is used to bring existing resources that were created outside of Terraform into the Terraform state. This allows Terraform to manage and track these resources as part of your infrastructure configuration.

- `terraform init` initializes the Terraform working directory and configures the backend but does not import resources into state.
- `terraform plan` creates an execution plan, showing what actions Terraform will take based on the current state and configuration, but it does not import resources.
- `terraform refresh` updates the state file to reflect the actual state of resources but does not add resources to the state if they are not already tracked by Terraform.

Therefore, to add existing resources into Terraform state, `terraform import` is the correct command.

Question 55**Correct**

The `.terraform.lock.hcl` file tracks module versions.

True

Explanation

False. The `.terraform.lock.hcl` file does not track module versions. It is used to lock the versions of the provider plugins used in a Terraform configuration to ensure consistency and reproducibility of the infrastructure deployment.

Your answer is correct

False

Explanation

True. The `.terraform.lock.hcl` file is not used to track module versions. Instead, it is specifically used to lock the versions of the provider plugins to maintain consistency and prevent unexpected changes in the infrastructure deployment process.

Overall explanation

Correct Answer:

False

Explanation:

The `.terraform.lock.hcl` file does **not track module versions**. Its primary purpose is to lock the versions of **providers** used in a Terraform configuration to ensure consistent and repeatable deployments across environments. Module versions, on the other hand, are not included in this lock file.

Key Details:

1. **What `.terraform.lock.hcl` Tracks:**
 - Provider versions and their checksums.
 - Specific versions of providers downloaded from registries.
 - This ensures all team members use the same provider versions, even if a newer version is available.
2. **What It Does Not Track:**
 - Terraform modules referenced in the configuration.
 - Module versions are managed directly in the Terraform code, typically via `source` and `version` attributes in the `module` block.
3. **Module Version Management:**
 - Terraform downloads modules separately and caches them in the `.terraform/modules` directory.
 - To ensure specific module versions, you explicitly define the `version` constraint when sourcing modules from a registry.

Example:

Provider Lock Example in `.terraform.lock.hcl`:

```
provider "registry.terraform.io/hashicorp/aws" {
  version = "4.5.0"
  hashes = [
    "h1:abcdef1234567890...",
    "h2:0987654321fedcba...",
  ]
}
```

Module Version Example in Code:

```
module "example" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "3.5.0"
}
```

Why "True" Is Incorrect:

Modules are not locked in `.terraform.lock.hcl`. Instead, they are directly managed through version constraints in the configuration, allowing developers to control module versions without relying on the lock file.

Resources

[Dependency Lock File](#)

Question 56Correct

Which provisioner invokes a process on the machine running Terraform?

`null-exec`

Explanation

The `null-exec` provisioner is a placeholder that does nothing. It is used when you do not want to run any commands or processes during provisioning. It does not invoke a process on the machine running Terraform, so it is not the correct choice for this scenario.

`remote-exec`

Explanation

The `remote-exec` provisioner is used to invoke a process on a remote machine after the resource is created. It does not run the process on the machine running Terraform, so it is not the correct choice for invoking a process on the local machine.

`file`

Explanation

The `file` provisioner is used to copy files or directories from the machine running Terraform to the resource being created. It does not invoke a process on the machine running Terraform, so it is not the correct choice for this scenario.

Your answer is correct

local-exec

Explanation

The local-exec provisioner is used to invoke a process on the machine running Terraform. It allows you to run commands locally on the machine where Terraform is being executed, making it the correct choice for invoking a process on the local machine.

Overall explanation

Correct Answer: local-exec

Explanation:

The **local-exec** provisioner runs a command on the machine where Terraform is executed, which is typically your local machine. This provisioner is useful for invoking processes or scripts locally after Terraform has completed resource creation or modification.

- **remote-exec** runs commands on the remote machine, not the machine where Terraform is executed.
- **file** is used to copy files from the machine running Terraform to a remote machine.
- **null-exec** is not a standard provisioner in Terraform. It seems to be a mistaken option, as there is no such provisioner named "null-exec" in Terraform documentation.

So, **local-exec** is the correct provisioner that runs a process on the machine where Terraform is running.

Resources

[local-exec Provisioner](#)

Question 57Correct

Define the purpose of state in Terraform.

State stores variables and lets you quickly reuse existing code

Explanation

State in Terraform is not used to store variables or reuse existing code. It is used to track the state of infrastructure resources and manage their configuration.

State codifies the dependencies of related resources

Explanation

State in Terraform does codify dependencies between resources, but this is not its sole purpose. It also tracks the current state of resources and their configurations, allowing Terraform to plan and apply changes effectively.

State lets you enforce resource configurations that relate to compliance policies

Explanation

While state in Terraform can help ensure that resource configurations are in compliance with policies, its primary purpose is not to enforce compliance policies. State primarily tracks the state of resources and their configurations.

Your answer is correct

State maps real world resources to your configuration and keeps track of metadata

Explanation

The correct choice. State in Terraform serves the purpose of mapping real-world resources to your configuration and keeping track of metadata such as resource IDs, attributes, and dependencies. This allows Terraform to manage infrastructure resources effectively and accurately.

Overall explanation

Correct Answer:

State maps real-world resources to your configuration and keeps track of metadata

Explanation:

In Terraform, the **state** file serves as a critical component for managing infrastructure. It performs the following functions:

- **Maps real-world resources to Terraform configurations:** Terraform uses the state file to keep track of resources that it has created or manages in the cloud or other systems. This allows Terraform to understand which resources exist and how they relate to your configuration.
- **Tracks metadata and dependencies:** The state file includes metadata about resources (e.g., IDs, attributes) and dependency information. This is necessary to ensure resources are managed correctly and in the right order during operations like `terraform apply`.

- **Enables efficient updates:** Terraform uses the state file to calculate the difference between the current state of infrastructure and the desired configuration, enabling it to apply only the necessary changes.

Why the Other Options Are Incorrect:

- **"State stores variables and lets you quickly reuse existing code":** Variables are not stored in the state file. Variables are typically defined in `.tf` files or passed at runtime. The state file focuses on tracking resources, not variables or code reuse.
- **"State lets you enforce resource configurations that relate to compliance policies":** State doesn't enforce policies directly. Compliance enforcement is done through tools like Sentinel or by designing the configuration itself.
- **"State codifies the dependencies of related resources":** While the state tracks dependencies to ensure resources are applied in the correct order, this is not the sole purpose of the state file. The primary role is mapping resources to configurations and tracking metadata.

Resources